

Threads

```
print_mul (int arr[], int len) {  
    int mul = 1, i;  
    for (i = 0; i < len; i++)  
        mul *= arr[i];  
    return mul;  
}  
  
print_sum (int arr[], int len) {  
    int sum = 0, i;  
    for (i = 0; i < len; i++)  
        sum += arr[i];  
    return sum;  
}
```

```
main ()  
{  
    int *arr1 = malloc (sizeof(int) * 100);  
    int arr2[100];  
    read_from_file (arr1, 100);  
    read_from_file (arr2, 100);  
    create_thread (print_mul, arr1, 100);  
    print_sum (arr2, 100);  
    return 0;  
}
```

Schedule

```
struct list *ready_list;
```

```
struct thread {  
    void *esp;  
};
```

- idle_thread

Schedule

```
struct list *ready_list;  
struct thread {  
    void *esp;  
};
```

```
create_thread () {  
    struct thread *t = malloc ();  
    t->esp = malloc(4096) + 4096;  
    enqueue (ready_list, t);  
}
```


Schedule

```

schedule (struct thread *cur_thread) {
    enqueue (ready_list, cur_thread);
    next_thread = dequeue (ready_list);
    push %eax
    ....
    push %edi
    cur_thread->esp = cur_esp;
    new thread. esp = next_thread->esp;
    pop %edi
    ...
    pop %eax
}


```

- Why schedule() doesn't save the current program pointer?



$cur_thread \rightarrow esp = esp;$

$esp = next_thread \rightarrow esp$



foo {
 schedule;
}
bar {
 schedule;
}

Race condition

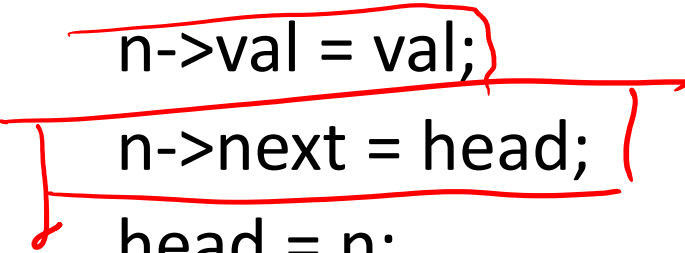
```
struct node *head; = NULL;
```

```
insert (int val) {  
    struct node *n = malloc();  
    n->val = val;  
    n->next = head;  
    head = n;  
}
```

Race condition

```
struct node *head = NULL;

insert (int val) {
    struct node *n = malloc();
    n->val = val;
    n->next = head;
    head = n;
}
```



```
thread 1:
n->next = NULL;
schedule ();
thread 2:
n->next = NULL;
head = n;
schedule ();
thread 1:
head = n;
```

Critical section

```
struct node *head = NULL;

insert (int val) {
    struct node *n = malloc();
    n->val = val;
    lock_acquire ();
    n->next = head;
    head = n;
    lock_release ();
}
```

The area between lock_acquire and lock_release is called a critical section.

Only one thread is allowed to execute in a critical section.

Thread 1.

lock_acquire

n->next = NULL;

Schedule

lock_acquire();

Schedule();

Thread 1: head = n;

lock_release();

Schedule();

Thread 2:

n->next = head;

head = n;

Implementation

```
lock_acquire ()  
{  
    status = interrupt_disable ();  
}
```

```
lock_release (int status)  
{  
    set_interrupt_status (status);  
}
```

insertC

Implementation

```
struct lock {  
    int value;  
};
```

```
lock_acquire (struct lock *l)  
{  
    while (l->value == 0) {}  
    l->value = 1;  
}
```

Thread 1:
l->value == 0
→ scheduloc()
l->value == 0

Thread 1:
l->value = 0;
→ l->value = 1;

```
lock_init (struct lock *l)  
{  
    l->value = 1 ;  
}
```

```
lock_release (struct lock *l)  
{  
    l->value = 0;  
}
```

Thread 2:
while (l->value == 0),
→ l->value = 0;

Implementation

```
struct lock {  
    int value;  
};
```

```
Lock_acquire (struct lock *l)  
{  
    while (l->value == 0) {}  
    l->value = 0;  
}
```

Race condition in the lock
implementation itself

```
Lock_release (struct lock *l)  
{  
    l->value = 1;  
}
```

Implementation

```
struct lock {  
    int value;  
};
```

```
Lock_acquire (struct lock *l)  
{  
    status = interrupt_disable();  
    while (l->value == 0) {}  
    l->value = 0;  
    set_interrupt_status (status);  
}
```

Race condition in the lock implementation itself

```
Lock_release (struct lock *l)  
{  
    l->value = 1;  
}
```

Implementation

```
struct lock {  
    int value;  
};
```

```
Lock_acquire (struct lock *l)  
{  
    status = interrupt_disable();  
    while (l->value == 0) {}  
    l->value = 1;  
    set_interrupt_status (status);  
}
```

Infinite loop:

```
Lock_release (struct lock *l)  
{  
    l->value = 0;  
}
```

Implementation

```
struct lock {  
    int value;  
};
```

```
Lock_acquire (struct lock *l)  
{  
    status = interrupt_disable();  
    while (l->value == 0) {  
        schedule ();  
    }  
    l->value = 1;  
    set_interrupt_status (status);  
}
```

```
Lock_release (struct lock *l)  
{  
    l->value = 0;  
}
```

Implementation

```
struct list *wait_list;
```

```
Lock_acquire (struct lock *l)
```

```
{  
    status = interrupt_disable();  
    while (l->value == 0) {  
        remove (ready_list, thread_current());  
        enqueue (wait_list, thread_current());  
        schedule ();  
    }  
    l->value = 0;  
    set_interrupt_status (status);  
}
```

```
Lock_release (struct lock *l)
```

```
{  
    struct thread *t = dequeue (wait_list);  
    enqueue (ready_list, t);  
    l->value = 1;  
}
```

Semaphore

```
struct list *wait_list;
```

```
sema_down (struct lock *l)
{
    status = interrupt_disable();
    while (l->value == 0) {
        remove (ready_list, thread_current());
        enqueue (wait_list, thread_current());
        schedule ();
    }
    l->value--;
    set_interrupt_status (status);
}
```

```
sema_up (struct lock *l)
{
    struct thread *t = dequeue (wait_list);
    enqueue (ready_list, t);
    l->value++;
}
```


Semaphore

```
parent () {  
    struct semaphore s;  
    sema_init (&s, 0);  
    create_thread (child, &s);  
    sema_down(&s);  
}  
child (struct semaphore *s) {  
    sema_up (s);  
}
```


```
char buf[BUF_SIZE];
size_t n = 0;
size_t head = 0, tail = 0;
struct lock lock;
struct condition not_empty, not_full;


void put (char ch) {
    while (n == BUF_SIZE) {}
    buf[head++ % BUF_SIZE] = ch;
    n++;
}
```

Condition variables

```
char get (void) {
    char ch;
    while (n == 0) {}
    ch = buf[tail++ % BUF_SIZE];
    n--;
}
```

```
char buf[BUF_SIZE];  
size_t n = 0;  
size_t head = 0, tail = 0;  
struct lock lock;  
struct condition not_empty, not_full;
```

```
void put (char ch) {  
    while (n == BUF_SIZE)  
        cond_wait (&not_full);   
    buf[head++ % BUF_SIZE] = ch;  
    n++;  
    cond_signal (&not_empty);  
}
```

```
char get (void) {  
    char ch;  
    while (n == 0)  
        cond_wait (&not_empty);  
    ch = buf[tail++ % BUF_SIZE];  
    n--;  
    cond_signal (&not_full);   
}
```

```
char buf[BUF_SIZE];
size_t n = 0;
size_t head = 0, tail = 0;
struct lock lock;
struct condition not_empty, not_full;
```

```
void put (char ch) {
    lock_acquire (&lock);
    while (n == BUF_SIZE)
        cond_wait (&not_full);
    buf[head++ % BUF_SIZE] = ch;
    n++;
    cond_signal (&not_empty);
    lock_release (&lock);
}
```

```
char get (void) {
    char ch;
    lock_acquire (&lock);
    while (n == 0)
        cond_wait (&not_empty);
    ch = buf[tail++ % BUF_SIZE];
    n--;
    cond_signal (&not_full);
    lock_release (&lock);
}
```

```
char buf[BUF_SIZE];
size_t n = 0;
size_t head = 0, tail = 0;
struct lock lock;
struct condition not_empty, not_full;
```

```
void put (char ch) {
    lock_acquire (&lock);
    while (n == BUF_SIZE) {
        lock_release (&lock);
        cond_wait (&not_full);
        lock_acquire (&lock);
    }
    buf[head++ % BUF_SIZE] = ch;
    n++;
    cond_signal (&not_empty);
    lock_release (&lock);
}
```

Thread 1



```
char get (void) {
    char ch;
    ✓ lock_acquire (&lock);
    while (n == 0) {
        lock_release (&lock);
        cond_wait (&not_empty);
        lock_acquire (&lock);
    }
    ch = buf[tail++ % BUF_SIZE];
    n--;
    ✓ cond_signal (&not_full); ✓
    lock_release (&lock);
}
```

```
char buf[BUF_SIZE];
size_t n = 0;
size_t head = 0, tail = 0;
struct lock lock;
struct condition not_empty, not_full;
```

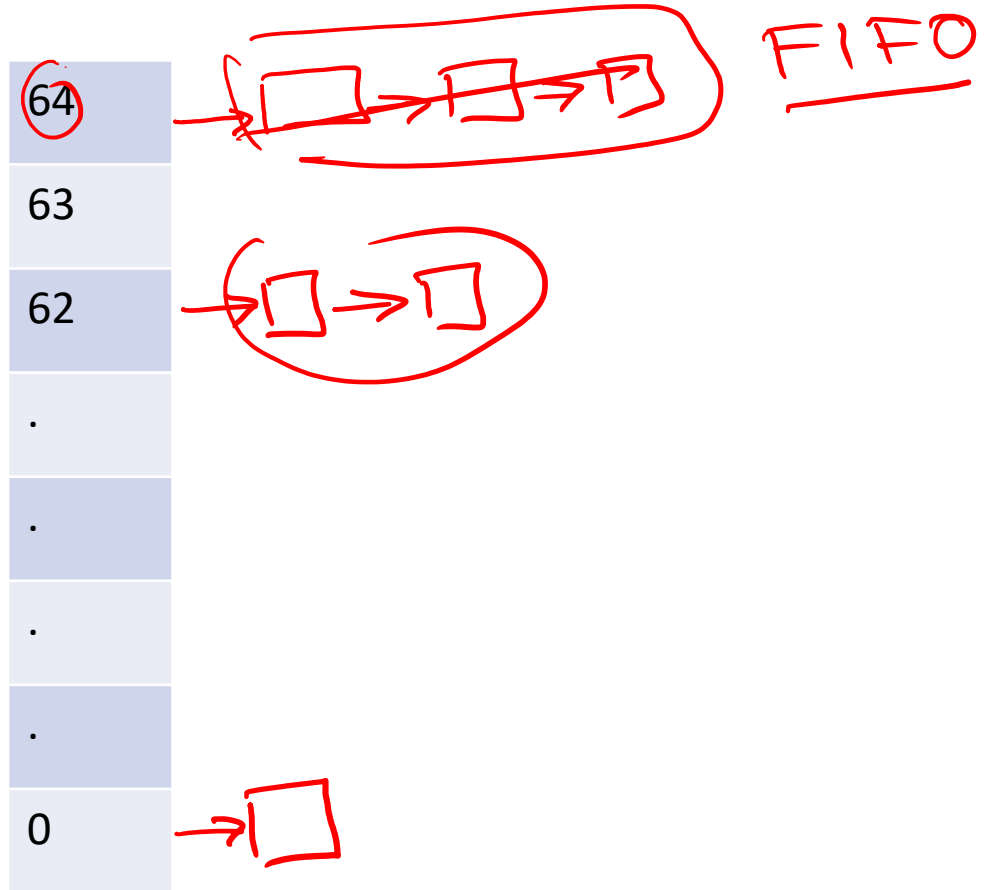
```
void put (char ch) {
    lock_acquire (&lock);
    while (n == BUF_SIZE)
        cond_wait (&not_full, &lock);
    buf[head++ % BUF_SIZE] = ch;
    n++;
    cond_signal (&not_empty);
    lock_release (&lock);
}
```

```
char get (void) {
    char ch;
    lock_acquire (&lock);
    while (n == 0)
        cond_wait (&not_empty, &lock);
    ch = buf[tail++ % BUF_SIZE];
    n--;
    cond_signal (&not_full);
    lock_release (&lock);
}
```

Priority scheduling

- Each thread is created with an initial priority
 - Priority is an integer value
- Let us say the priority range is 0-64
 - A higher value means a high priority
- In round-robin scheduling, the threads are scheduled in a FIFO order
- In priority scheduling threads with the same priority are scheduled in a FIFO order
 - The scheduler always schedules a high priority thread

Priority scheduling



Starvation

- In priority scheduling, a low priority thread may wait infinitely to get scheduled
- To prevent this, an OS periodically (say 15minutes) increase the priority of a low priority thread
 - also called aging
- After a low priority thread (with enhanced priority) gets scheduled the OS sets the priority to its original value

Assignment-1

- Priority scheduling
- Priority donation
- Deadline 9th Feb (Hard deadline, no extensions)
- Design document

