




# Processes

- Processes are containers of threads
- Each process start with one thread and can create more threads if needed
- Every process has its own quota in RAM
- A process cannot see other processes data



# Application

- OS is a shared library
- The library exports some interfaces that application developers can call to use OS services
- The OS enforces that an application can only call these exported routines

# Application


- Application is partitioned into two components
  - Application code (user programs)
  - OS library
- User programs are untrusted
- OS library is trusted

# How user programs use the OS library?

```
create_thread () {  
    struct thread *t = malloc (sizeof(struct thread));  
    t->esp = malloc (4096) + 4096;  
    status = interrupt_disable ();  
    add_to_ready_list (t);  
    set_interrupt_status (status);  
}
```


# Memory map

- User program and OS (kernel) lives in different address spaces
- Users programs cannot directly access kernel memory
- However, the kernel can access the entire memory



# Interrupt handling

$esp = Ustack + 100$




$mov \$8, -1(esp)$

iret



$mov \$8, -1(esp)$





# Why kernel stack is needed

# Software interrupts

128

- int \$100
- Linux uses vector 128 for system call
- The syscall id is passed in some registers or user stack

```
idt[128] = syscall_handler  
mov $0, %eax  
int $128
```

```
syscall_handler(int id)  
{  
    switch(id) {  
        case 0:  
            create_thread();  
            break;  
        case 1:  
            create_process();  
            break;  
    }  
}
```


# Unix operating system

- “shell” is the first user program created by the Unix
- “shell” can create more processes

# Fork


```
main () {  
    int pid;  
    pid = fork();  
    if (pid == 0) {  
        /* child process */  
    } else if (pid > 0) {  
        /* parent process */  
    }  
}
```

200 ↑ Parent  
child  
0



# Exec

```
main (int argc, char *argv[]) {  
    exec ("ls", argv, 0);  
}
```



# System calls

- `int creat (pathname, mode)`
- `write (fd, buf, len)`
- `read (fd, buf, len)`
- `close (fd)`

```
fd = creat ("tmp.txt", 0666)
write (fd, "Hello world", sizeof("Hello..."));
char buf[64];
len = read (fd, buf, 64);
close (fd);
```

# System calls

- fds 0, 1, 2 have special meaning
- 0 points to standard input
  - e.g., keyboard
- 1 points to standard output
  - e.g., terminal
- 2 points to standard error
  - e.g., terminal

# Shell

```
while (1) {  
    write (1, "$ ", 2);  
    readcommand (0, command, args);  
    if ((pid = fork ()) == 0) {  
        exec (command, args, 0);  
    } else if (pid > 0) {  
        wait (0);  
    } else  
        printf ("Failed to fork\n");  
}
```



# exit (int status)

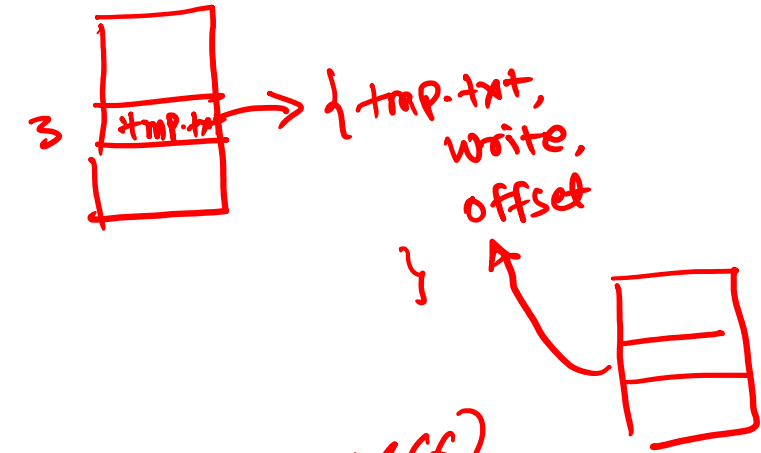
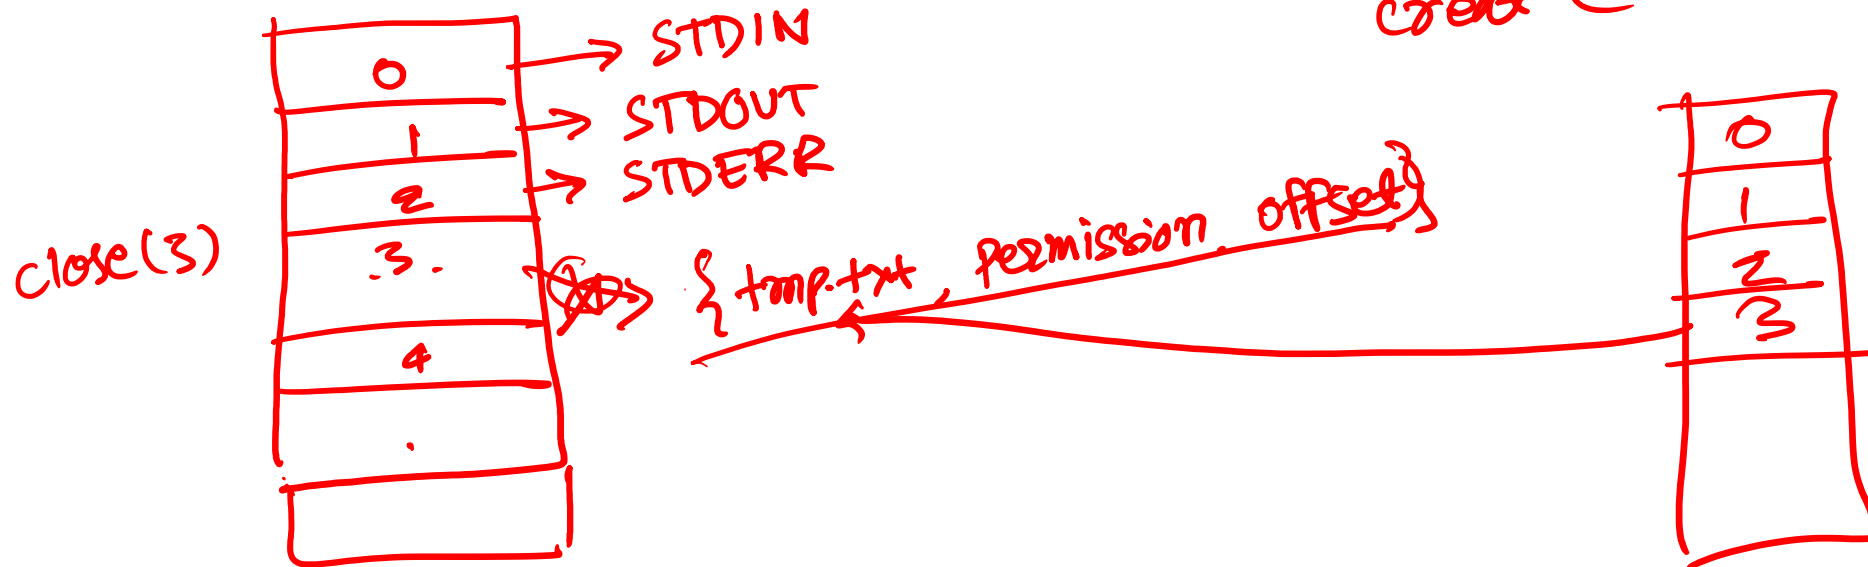
- The process terminates with a given status
- Wipe out all the memory occupied by the process
- Release all resources

# int wait ()

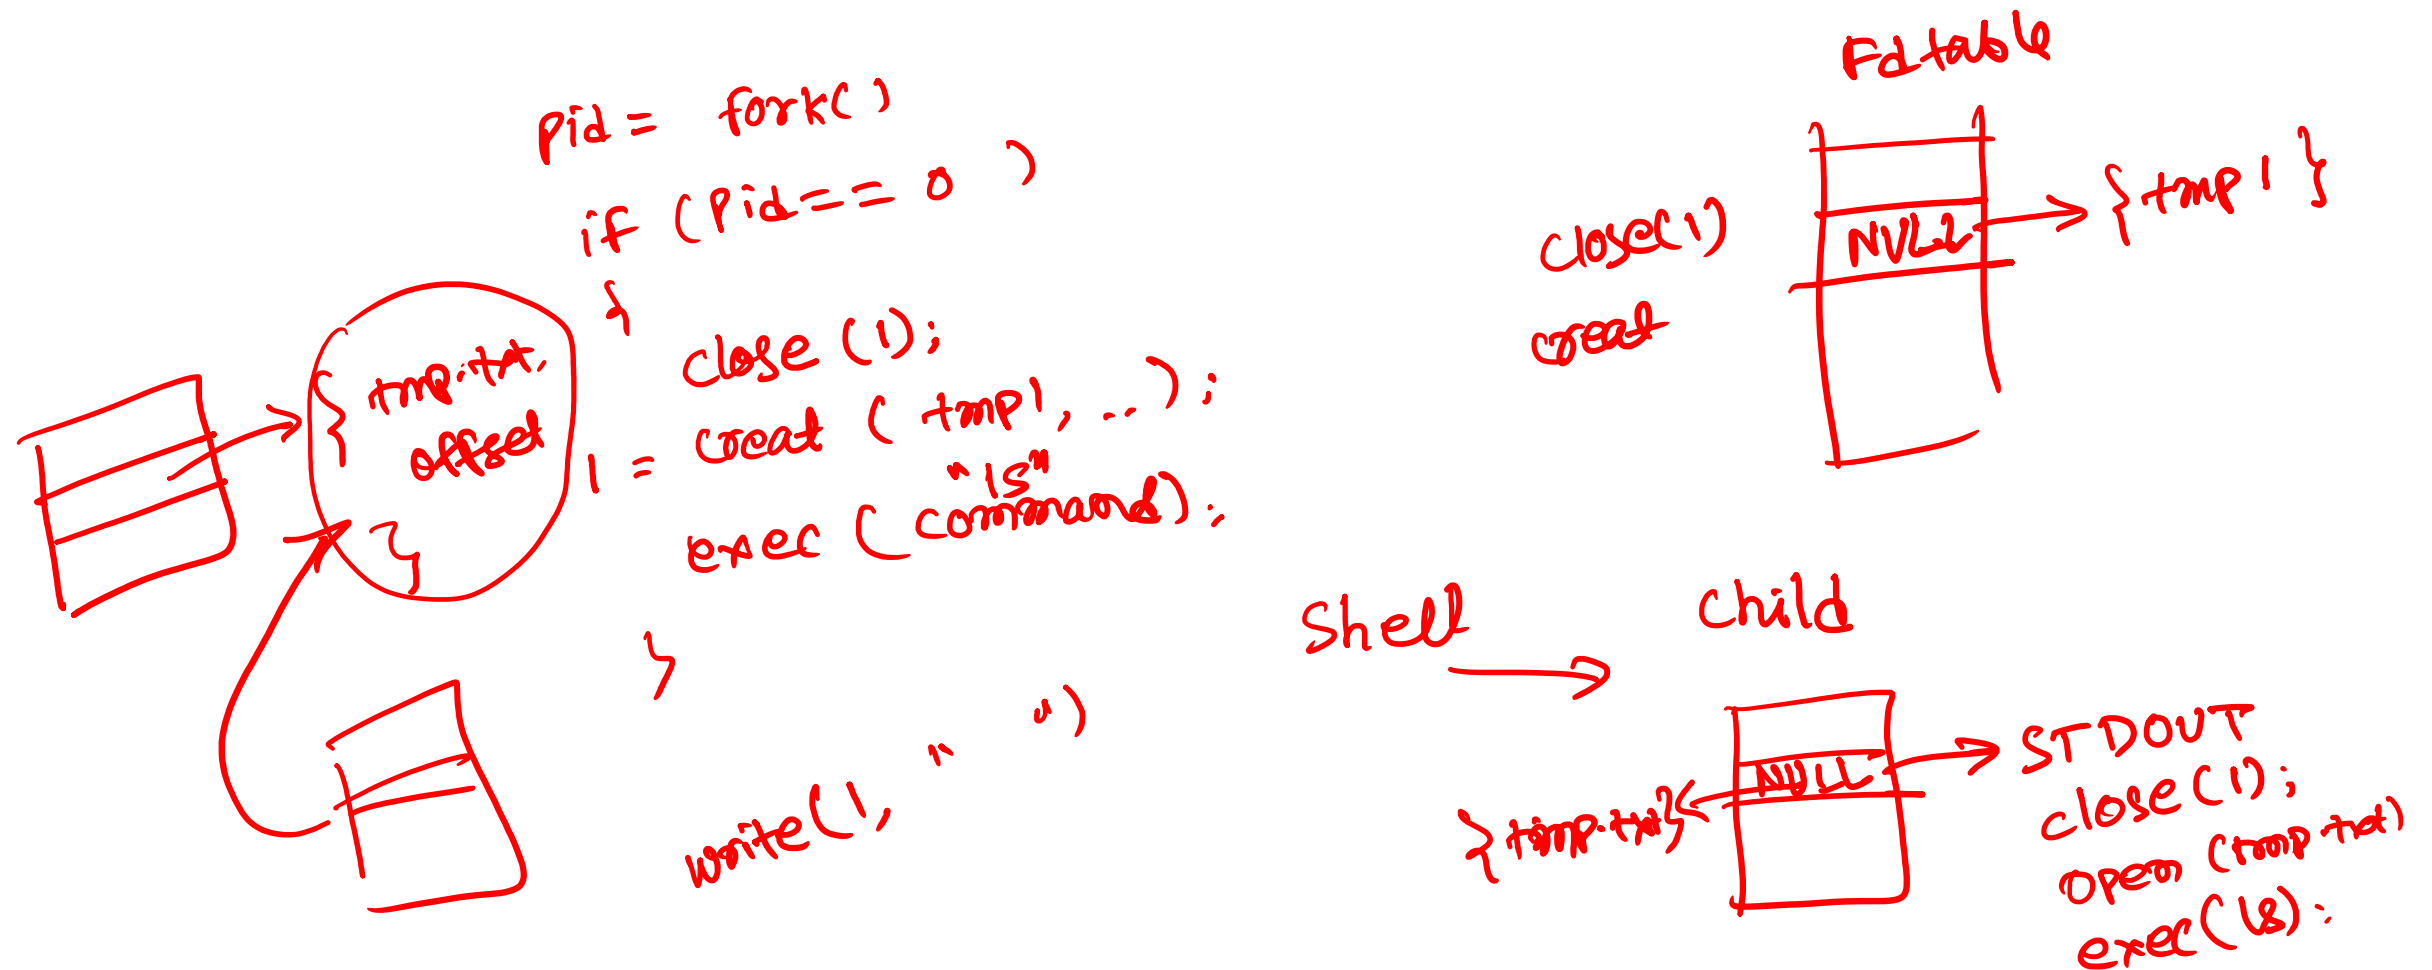
- wait system call waits until the child terminates
- The return value is the exit status of the child

# File descriptor table

- File descriptors are inherited by the child process



# How does shell implement "ls > tmp1"



sh < script > tmp1

<      close (0);  
         open (script);  
     >      close (1);  
             creat ("tmp1");

sh

read\_command (0,

1 → tmp1

Is f1 f2 nonexistent-f3 > tmp1 2>&1

STDERR

↓  
(2)


$f_{d1} = \text{open}(\text{tmp.txt})$   
 $f_{d2} = \text{open}(\text{tmp.txt})$

close(1);  
creat(tmp1);

~~close(2)~~  
creat(tmp1);


dup(1);

close(1);  
creat(tmp1);  
close(2);  
dup(1);



{ tmp1.txt;  
write;  
offset=0 } → 100

{ tmp1.txt,  
write  
offset=0




tmp1

dup system call

# Inter process communication

- Pipe
  - A pipe has two end
  - Data are written to the input end
  - Data are fetched from the output end






# Pipes


```
int fd[2];  
char buf[512];  
int n;
```

```
pipe (fd);  
write (fd[1], "hello", 5);  
n = read (fd[0], buf, sizeof(buf));  
// buf[] now contains 'h', 'e', 'l', 'l', 'o'
```



# Inter process communication

```
int fd[2];  
char buf[512];  
int n, pid;  
  
pipe (fd);  
pid = fork ();  
if (pid > 0) {  
    write (fd[1], "hello", 5);  
} else {  
    n = read(fd[0], buf, sizeof (buf));  
}
```



# How to run a series of programs?



```
sort < file.txt > tmp1
```

```
uniq < tmp1 > tmp2
```

```
wc < tmp2
```

```
rm tmp1 tmp2
```

file.txt




sort < file.txt | uniq | wc

# How to run a series of programs?



```
sort < file.txt | uniq | wc
```

# Pipes



```

pipe (fd) ; ✓
if ((pid == fork()) == 0) {
    close (1);
    tmp = dup (fd[1]);
    close (fd[0]);
    close (fd[1]);
    exec (command1sort, args1, 0);
} else if (pid > 0) {
    close (0);
    tmp = dup (fd[0]);
    close (fd[0]);
    close (fd[1]);
    exec (command2uniq, args2, 0); }
    
```



```

fork ( ) ;
if (pid >= 0) {
    exec ("sort") ;
}
else {
    exec ("uniq") ;
}
    
```