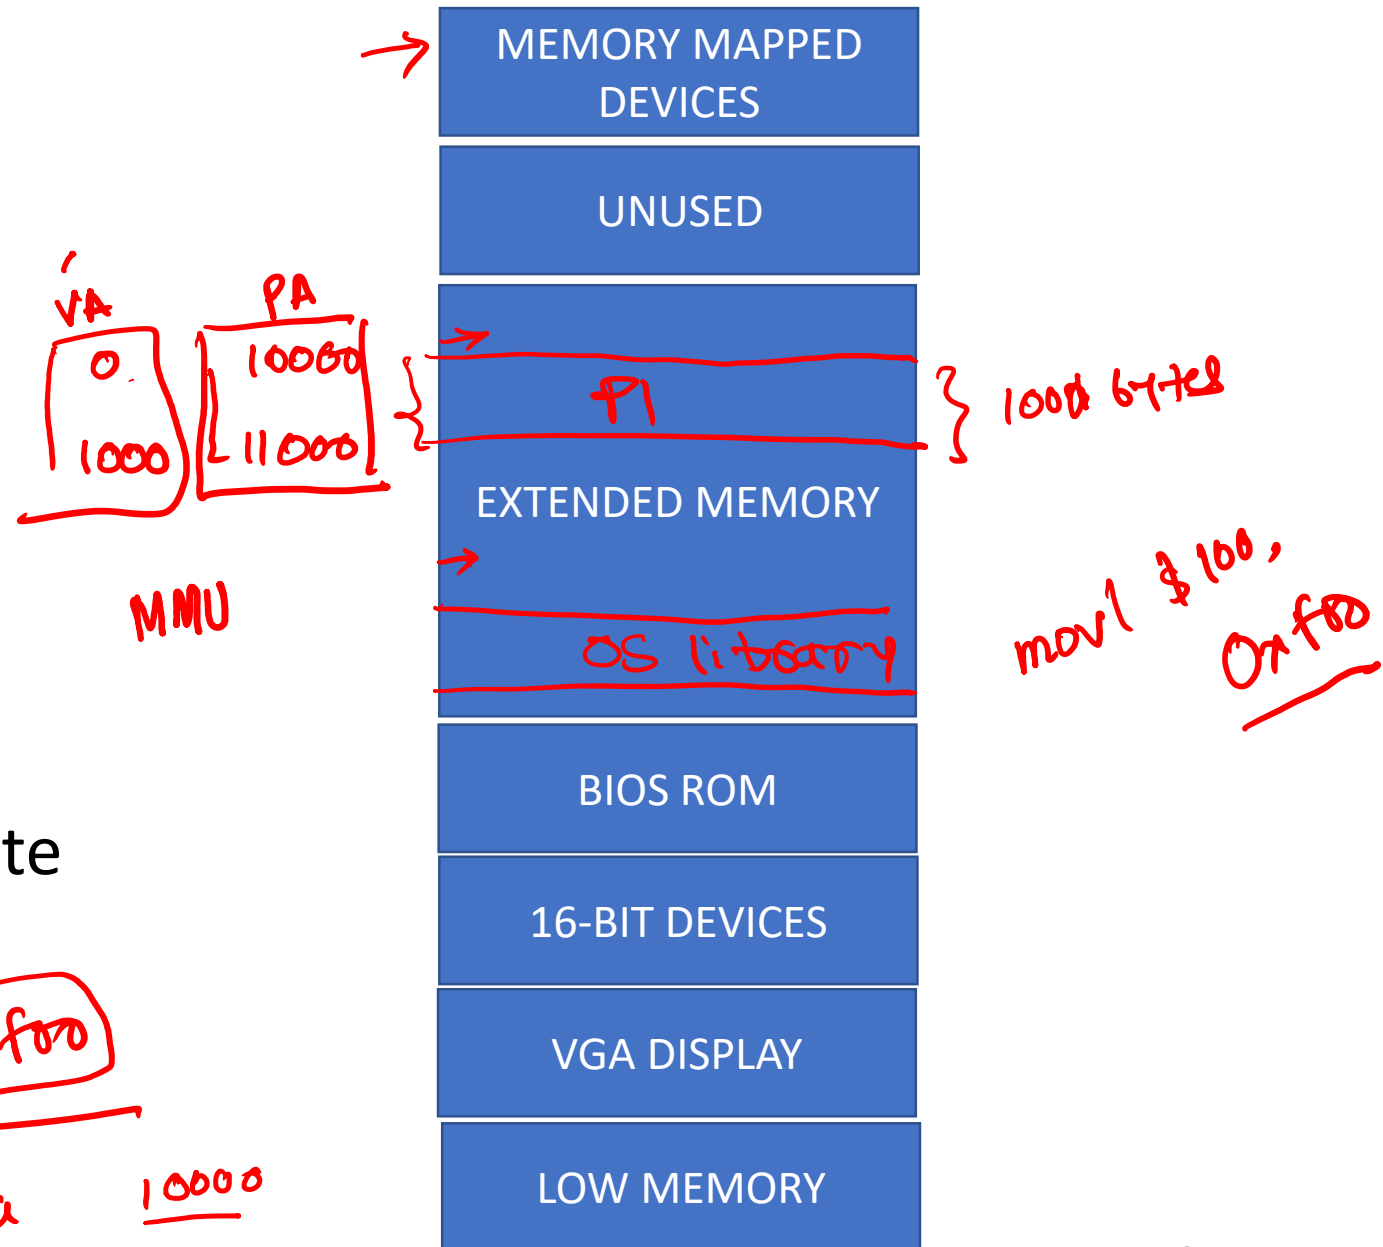


Memory isolation

- Processes have their own reserved quota of RAM
 - also called the address space of the process
- Decided by OS during fork/exec


Memory map

- The compiler does not know actual RAM addresses in advance
- If we only allow physical addresses, the compiler is unnecessarily forced to generate relocatable code



Virtual addresses

- Uniform among all processes
- Converted into physical address (PA) by hardware during memory access
- Virtual address (VA) range is 0 to size of process reserved quota in RAM
- PA is computed using adding the base of the process address space to VA



Segmentation

- To facilitate address translation CPU has an additional unit called MMU
- In addition to general purpose registers, x86 also have six segment registers namely, %cs, %ds, %ss, %es, %fs, and %gs

Segmentation

- Every memory operand is prefixed by some segment register
- Let us assume for now that %cs segment register is used for all memory operands


1. eax
movl \$100, -1ds(-1. eax)
mov \$100, -1fs(-1. eax)
1. cs(-1. eax)

Segmentation



- Segment registers contain a 16-bit value that is used to index in a global descriptor table (GDT)
- GDT is an array of structures of the following type.

```
struct {  
    unsigned base, limit, perm;  
};
```

$\cdot 1 \cdot CS : (\cdot 1 \cdot eax)$
address of eax



Global descriptor table



\uparrow
 $\text{mov } \$103, 10102$

Segmentation

- The MMU unit adds the base of the descriptor to the virtual address to compute the physical address during a memory access
- If the virtual address exceeds a limit, the hardware generates an error

```
movl $100, %cs:100
```

```
index = cs >> 3; // cs is 16-bit
```

```
base = gdt[index].base;
```

```
unsigned *phys = base + 100;
```


```
assert (100 < gdt[index].limit);
```

```
*phys = 100
```


GDT

```
mov 0, %ax
mov %ax, %ds
movl $1, %ds:100
%ds:100 = 1000 + 100
assert (100 < 10000)
```

```
mov 16, %ax
mov %ax, %ds
movl $1, %ds:100
%ds:100 = 200 + 100
assert (100 < 256)
```



	base	limit	system/supervisor
0	1000	10000	
1	30000	31000	
2	200 0	256 0xFFFF	
3	
4	
5	
6	
7	

RAM


1000
16000
P1

Who creates GDT?

- OS creates GDT
- OS creates GDT in RAM and executes “lgdt” instruction to load the address of GDT in GDTR register
- Why can't user programs do the same thing?

Protection rings

- X86 has 4 protection rings (0,1,2, and 3)
- 0 is most privileged and, 3 is the least privileged mode
- OS executes in most privileged mode
- User programs execute in the least privileged mode



Protection rings

- Some instructions are only allowed in most privileged mode
- Some instructions have different semantics in different privileged mode
- “lgdt” is only allowed in most privileged mode

How does hardware know the current privilege level (CPL)?

- The last two bits of ^{cs} segment register contains the CPL
 - For OS CPL == 0
 - For user programs CPL == 3

What prevent applications from modifying the GDT entries?

- GDT lives in OS memory

How does OS access the entire memory?

- OS segment selectors point to a GDT entry with base and limit set to 0 and 2^{32} respectively
- Why user programs cannot set their segment registers to OS GDT entries
 - Supervisor (system) flag in the GDT entry

ds - 0

ds - (2<<3)

1000	2000	0
0	0xffffffff	1

Context switch

- A dedicated GDT for every process
- Load the target process' GDT on every context switch
 - execute lgdt

p2 lgdt 7
 p1 lgdt 8

jmp .cs: eip

p1

X	0	1000	2000	0
8		3000	4000	0
✓		0	0xffffffff	1

p1 :
ds ← 0
p2 :
ds ← 8

3000

p2

7		3000	4000	0
✓				
✓		0	0xffffffff	1

Context switch

- Only one GDT
- Overwrite the GDT entries with target process entries on every context switch
 - modify GDT table itself

1000	2000	0
0	0xffffffff	1

Default segment registers

- %cs is default for EIP
- %ss is default of stack
 - e.g., push, pop
- %ds is default for most of memory operands
 - e.g., movl \$100, (%eax) /* default ds */
- %es is default for string instructions *.! es*
 - e.g., movsb
- %ds can be overridden by other segment registers
 - e.g., movl \$100, %fs(%eax) /* use fs instead of ds */

Exceptions

- Divide by zero
 - CPU executes divide by zero
- General protection fault
 - Process tries to access a VA outside its limit
- and many more ...
- 0 – 32 interrupt vectors are reserved for exceptions

Interrupt handler

- IDT is an array of structures of 256 entries
- IDT contains cs:eip pairs of interrupt handlers

```
struct {  
    unsigned interrupt_handler;  
    unsigned short cs;  
    ...  
};
```

Who creates IDT?

- OS creates IDT
- OS creates IDT in its own address space and executes “lidt” instruction to load the address of IDT in IDTR register
- Why can't user programs do the same thing?
 - “lidt” can only be executed in ring 0

Task state segment (TSS)

- On interrupt in user programs, the hardware automatically switch to the kernel stack
- A special structure TSS contains the address of kernel stack and kernel stack segment

Interrupt/exception user mode

switch to kernel stack (ss:esp)

push old_ss ✓

push old_esp ✓

push old_eflags ✓

push old_cs ✓

push old_eip ✓

Interrupt/exception in kernel mode

push old_eflags

push old_cs

push old_eip

Returning to user mode

iret

pop old_eip

pop old_cs

pop old_eflags

pop old_esp

pop old_ss

Returning to kernel mode

iret

pop old_eip

pop old_cs

pop old_eflags

Fork and exec

iret


pop old_eip

pop old_cs

pop old_eflags

pop old_esp

pop old_ss



Software interrupts

- int3

- Generates interrupt of vector number 3

GDB

main()
(mov .1.eax, .1.ecx)
int3



- into

- Generates interrupt of vector number 4, if the overflow condition is set

-- try {

}
catch {
}

call handler
5 bytes



divide_by_zero
{