

Process control block (PCB)

```
struct process {  
    int pid;  
    unsigned base;  
    unsigned limit;  
    struct fd *fdtable;  
    struct list threads;  
};
```

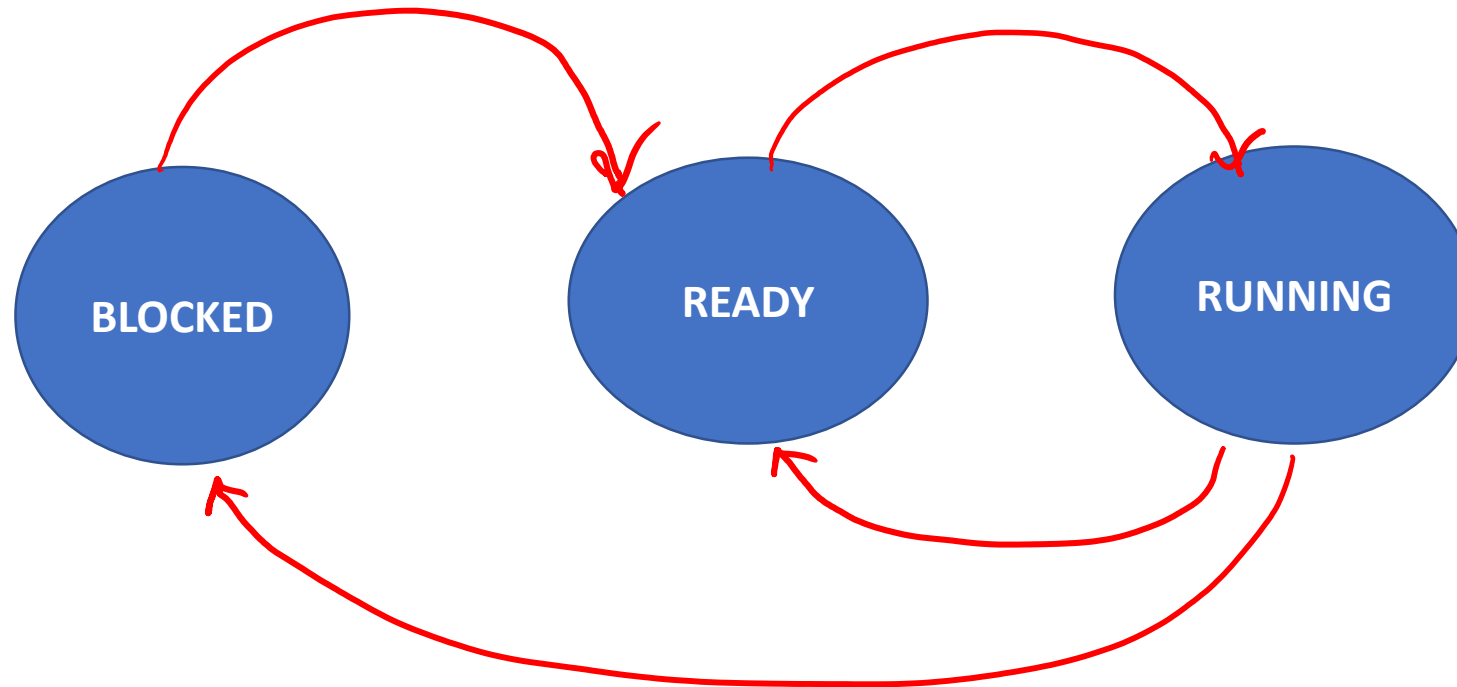
Thread control block (TCB)

```
struct thread {  
    int tid;  
    struct process *pcb;  
    void *esp;  
    unsigned regs[8];    // general purpose registers  
    int state;  
};
```

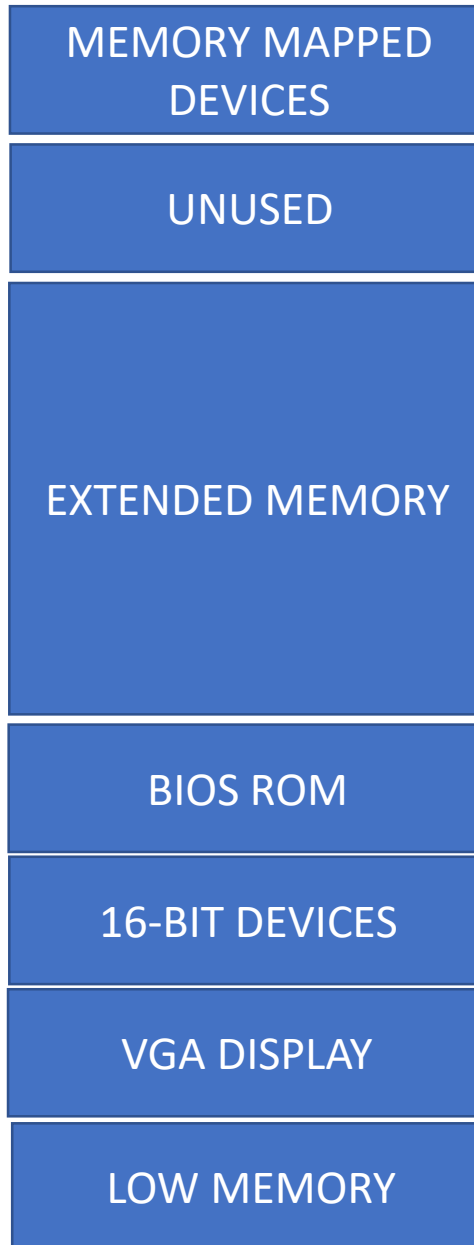
Threads

- thread state can be one of these three states
 - RUNNING, READY, BLOCKED

Threads



Memory map




Port I/O

- Use dedicated I/O space
 - Only 1024 I/O address
 - accessed using in/out instructions

in \$100

out \$100



MEMORY MAPPED
DEVICES

UNUSED

EXTENDED MEMORY

BIOS ROM

16-BIT DEVICES

VGA DISPLAY

LOW MEMORY

Line printer (Port I/O)

```
#define DATA_PORT 0x378
```

```
#define STATUS_PORT 0x379
```

```
#define BUSY 0x80
```

```
#define CONTROL_PORT 0x37A
```

```
#define STROBE 0x01
```

```
void lpt_putc(int c) {
```

```
    while(((inb(STATUS_PORT) & BUSY) == 0) ; /* wait for printer to consume previous byte */
```

```
    outb(DATA_PORT, 'c'); /* put the byte on the parallel lines */
```

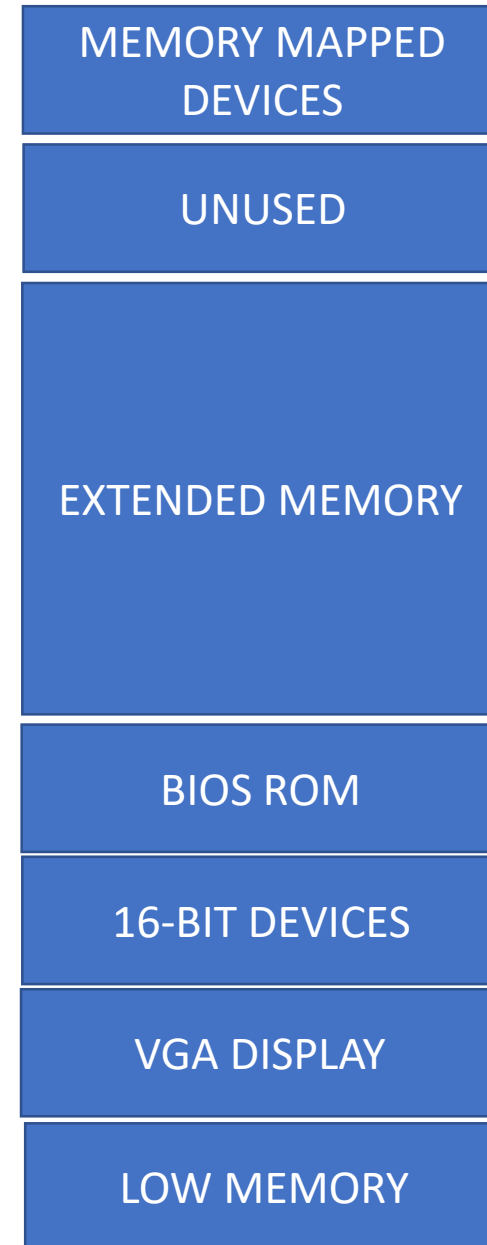
```
    outb(CONTROL_PORT, STROBE); /* tell the printer to look at the data */
```

```
    outb(CONTROL_PORT, 0);
```

```
}
```


Memory mapped I/O

- Some addresses in the physical address space are reserved for devices
- Software talks to the devices by reading/writing to these addresses



Line printer (MMIO)

```
volatile int *DATA_REG = 0xffffffff00;  
volatile int *STATUS_REG = 0xffffffff04;  
volatile int* CONTROL_REG = 0xffffffff0C;  
#define BUSY 0x80  
#define STROBE 0x01
```

```
void lpt_putc(int c) {  
    while ((STATUS_REG[0] & BUSY) == 0) ; /* wait for printer to consume previous byte */  
    DATA_REG[0] = 'c'; /* put the byte on the parallel lines */  
    CONTROL_REG[0] = STROBE; /* tell the printer to look at the data */  
    CONTROL_REGREG[0] = 0;  
}
```

Volatile vs non-volatile

```
while ((STATUS_REG[0] & BUSY) == 0) ;
```

```
int reg = STATUS_REG[0];  
while ((reg & BUSY) == 0);
```

Direct memory access (DMA)

```
#define TX_STATUS_PORT 0x400
#define TX_ADDR_PORT   0x401
#define RX_STATUS_PORT 0x402
#define RX_ADDR_PORT   0x403
```

```
tx_packet(char *buf) {           // transmit buf to the network
    outl(TX_STATUS_PORT, 0);
    outl(TX_ADDR_PORT, virt_to_phys(buf));
    while (inl(TX_STATUS_PORT) == 0); // polling until buf is sent
}
```

waiting

Direct memory access

```
#define TX_STATUS_PORT 0x400
#define TX_ADDR_PORT   0x401
#define RX_STATUS_PORT 0x402
#define RX_ADDR_PORT   0x403
```

```
rx_packet (char *buf) {           // receive network packet in buf
    outl (RX_STATUS_PORT, 0);
    outl (RX_ADDR_PORT, virt_to_phys(buf));
    while (inl(RX_STATUS_PORT) == 0); // wait until data is received in buf
}
```

Polling

- Periodically checking for the occurrence of an event
 - e.g. after every 100 ms network device checks for the arrival of a packet
- Polling is not useful in all cases
 - Specifically, when the device events are not frequent, e.g.,
 - the keyboard input
 - low throughput network channel

Interrupts

```
#define TX_STATUS_PORT 0x400
#define TX_ADDR_PORT 0x401
#define RX_STATUS_PORT 0x402
#define RX_ADDR_PORT 0x403

rx_packet (char *buf) {
    outl (RX_STATUS_PORT, 0);
    outl (RX_ADDR_PORT, virt_to_phys(buf));
    //while (inl(RX_STATUS_PORT) == 0);
}
```

```
/* called after packet is received */
rx_intr () {
    addr_t phys = inl(RX_ADDR_PORT);
    char *buf = phys_to_virt (phys);
    return buf;
}

/* called after every 100 ms */
rx_poll() {
    addr_t phys = inl(RX_ADDR_PORT);
    char *buf = phys_to_virt (phys);
    return buf;
}
```

Disk driver

- Read disk driver section from chapter 3 of xv6-book for a real device driver

Shell

- What if the shell doesn't call wait system call
- The shell can execute another command even if the previous command hasn't finished yet

```
while (1) {  
    write (1, "$ ", 2);  
    readcommand (0, command, args);  
    if ((pid = fork ()) == 0) {  
        exec (command, args, 0);  
    } else if (pid > 0) {  
        status = wait (0);  
    } else  
        printf ("Failed to fork\n");  
}
```

Background process in shell

- ls &
 - run "ls" in background
 - shell doesn't wait for ls to complete before reading another command

```
while (1) {  
    write (1, "$ ", 2);  
    readcommand (0, command, args);  
    if ((pid = fork ()) == 0) {  
        exec (command, args, 0);  
    } else if (pid > 0) {  
        //status = wait (0);  
    } else  
        printf ("Failed to fork\n");  
}
```

Problem with not calling wait

- The child process is not completely destroyed until the parent process calls the wait system call
- The wait system call returns the exit status of the child process
- The child process becomes a zombie process after doing an exit system call until the parent process does the wait system call

Zombie process

- A zombie process is a resource leak
- The shell somehow need to call the wait system call for the background process

Signals

- SIGINT (ctrl + c)
- SIGSTOP (ctrl + z)
- SIGCHLD (child process terminates)
- SIGTRAP (processor executed int3)
- SIGSEGV (memory access outside limit)
- and may more

Signal system call

```
sighandler_t signal(int signum, sighandler_t handler);
```

- Registers a handler for a given signal with the OS. The handler is called when the signal is generated.
- Signals invoke an asynchronous method in the user address space, similar to interrupts.

Signals


/* sigHandler is called every time "ctrl + c" is pressed */

```
void sigHandler(int signal) {  
    printf ("received %d signal\n", signal);  
}
```

```
main () {  
    signal (SIGINT, sigHandler);  
    ...  
}
```



user process



```
terminal_driver()  
{  
    SIGINT  
    sigHandler  
}
```

kill system call

- `int kill (pid_t pid, int sig)`
- Send a signal to another process
- e.g., a process can send SIGKILL signal to kill a target process


Background process in shell

```
childExit()  
{  
    wait (0);  
}
```

```
signal (SIGCHLD, childExit);  
while (1) {  
    write (1, "$ ", 2);  
    readcommand (0, command, args);  
    if ((pid = fork ()) == 0) {  
        exec (command, args, 0);  
    } else if (pid > 0) {  
        //status = wait (0);  
    } else  
        printf ("Failed to fork\n");  
}
```


Fragmentation

- Suppose the RAM of size 10000 bytes is mapped at location 0 – 10000
- And we have 10 processes (P1 - P10) each of size 1000 bytes




Fragmentation

- Let's say P1 and P10 have finished their execution
- At this point, the OS wants to load a new process P11 of size 2000
- The OS needs to relocate either P2 or P9 to accommodate the new process
 - need to copy the entire process



Fragmentation

- Similarly, if a process needs more RAM during its execution and consecutive addresses are not available the OS may need to relocate the entire process



Process address space in segmentation

- The process address space is limited by the amount of RAM available on the system

Paging

- Due to the limitations of segmentation the x86 hardware has an additional memory management hardware called paging hardware
- In paging scheme, the process virtual address space is $(0 - 2^{32}-1)$ for every process