


Executable and linkable format (ELF)

```
int numentries = 102400;  
int buf[102400];  
int main () {  
    read_from_file (buf, numentries);  
    sort (buf, numentries);  
    print_data (buf, numentries);  
    return 0;  
}
```



ELF

```
int numentries = 102400;  
int buf[102400];  
int main () {  
    read_from_file (buf, numentries);  
    sort (buf, numentries);  
    print_data (buf, numentries);  
    return 0;  
}
```

.text

cpu instructions of main,
read_from_file, sort, print_data,
etc.

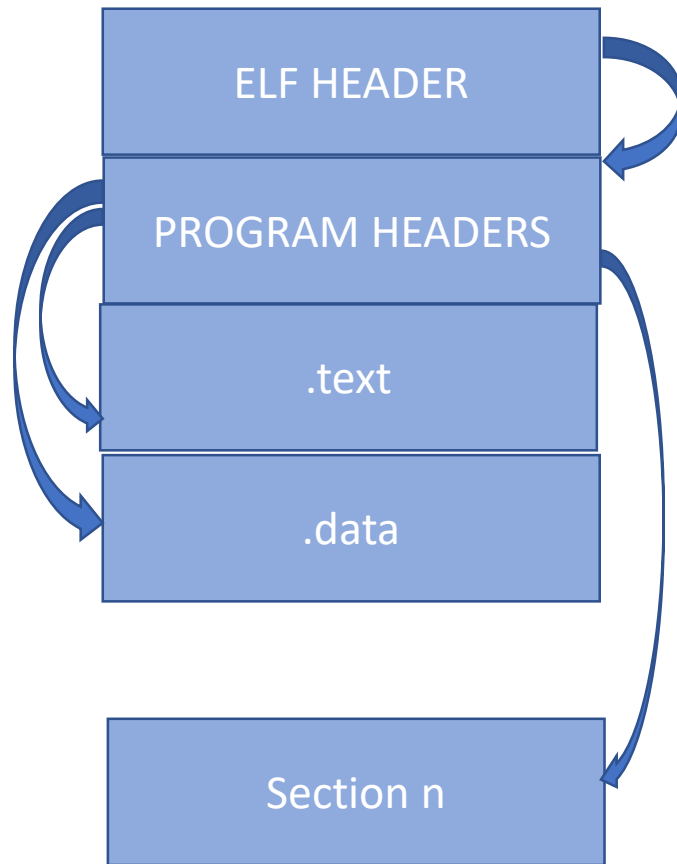
.data

numentries, buf

.stack

ELF


a.out




elfhdr: 1005

```
struct elfhdr {  
    uint magic;    /* must equals to 0x464C457FU */  
    uchar elf[12];  
    ushort type;  
    ...  
    uint phoff; /* offset of first program header */  
    ...  
    ushort phnum; /* number of program headers */  
}
```

proghdr: 1024

```
struct proghdr {  
    uint type;  
    uint off;  
    uint vaddr;   
    uint paddr;  
    uint filesz;  
    uint memsz;  
    uint flags;  
    uint align;  
};
```

Process address space



- text
- data
- unmapped
- stack
- heap

exec: 6310

`exec (char *path, char **argv)`


- Open executable file
- Read elf header from the executable file
- Create page directory using `setupkvm`

setupkvm:1837

setupkvm()

- Create page directory
- map I/O space
 - {"0", "1MB"} -> {"KERNBASE", "KERNBASE+1MB"}
- map kernel text
 - {"1MB", "textsize"} -> {"KERNBASE+1MB", "KERNBASE+1MB+textsize"}
- map rest of the RAM
 - {"1MB+textsize", "ramsize"} -> {"KERNBASE+1MB+textsize", "KERNBASE+ramsize"}

va 




mappages:1779

mappages (pgdir, va, size, pa, perm)

- Map physical address range {pa, pa+size} to virtual address range {va, va+size}

Virtual address space after setupkvm



exec: 6310


exec(char *path, char **argv)

vaddr ✓
memsz ✓

- read program headers from file
- for every program header
 - fetch the virtual address (vaddr) and memory size (memsz) of the section from the program header
- call allocvm to map physical pages at the virtual address range {vaddr, vaddr+memsz}


allocvm: 1953

allocvm (pgdir, oldsz, newsz)



- allocate physical pages for all the virtual pages between oldsz and newsz
- allocate physical pages using kalloc
- map physical pages using mappages
- return newsz on success

Virtual address space after allocuvm




exec: 6310

`exec(char *path, char **argv)`

- read program headers from file
- for every program header
 - fetch the virtual address (vaddr) and memory size (memsz) of the section from the program header
- call `allocvm` to map physical pages at the virtual address range {vaddr, vaddr+memsz}
- call `loadvm` to load the content of the section from the file


loadvm: 1918

loadvm (pgdir, addr, ip, offset, sz)



- read “sz” bytes from the input file at the given offset and store them at the virtual address “addr” mapped in the input pgdir
- Why does loadvm walk the page directory to first get the physical address instead of directly copying into the virtual address?
 - pgdir may not be the currently active page table

Virtual address space after loadvm




exec: 6310

`exec(char *path, char **argv)`

- read program headers from file
- **for every program header**
 - fetch the virtual address (vaddr) and memory size (memsz) of the section from the program header
 - call `allocvm` to map physical pages at the virtual address range {vaddr, vaddr+memsz}
 - call `loadvm` to load the contents of the section from the file


Virtual address space after for loop




exec: 6310

exec (char *path, char **argv)

- read program headers from file
- load all sections
- allocate two pages for stack
- make the first page inaccessible
 - by revoking the user access



Virtual address space after stack allocation



exec: 6310

`exec (char *path, char **argv)`

- read program headers from file
- load all sections
- allocate two pages for stack
- make the first page inaccessible
 - by revoking the user access
- **copy arguments of main to stack**

Copy arguments to stack

- Our goal is to call “main (int argc, char *argv[])” on this stack
- exec is called with null terminated “char **argv”
- for each string in argv
 - allocate space on stack
 - copy the string to the allocated space
- set up the arguments, a fake return address on the stack

Copy arguments to stack

```
char *argv[] = {"foo", "bar", NULL};
```

```
argc = 2;
```


exec (argv)

arg[] = {foo,
 bar
 };

for

base

```
main( int argc, char *argv[])
```



exec: 6310

`exec (char *path, char **argv)`

- read program headers from file
- load all sections
- allocate stack and copy arguments
- **adjust process size**

exec: 6310

`exec (char *path, char **argv)`

- read program headers from file
- load all sections
- allocate stack and copy arguments
- adjust process size
- rewrite eip with the main of executable and esp with the new stack in the trap frame (pushed during the exec system call)

trapframe : 602

```
struct trapframe {  
    uint edi;  
    ...  
    uint eax;  
    ushort gs;  
    ushort padding1;  
    ...  
    uint trapno;  
    uint err;  
    uint eip;  
    ushort cs;  
    ...  
};
```

alltraps: 3254

ss, esp, eflags, cs, eip // pushed by hardware
error_code, vector no // pushed by vectors.S

alltraps:

push all segment registers

push all general purpose registers

trap (&esp) /* esp contains the address of
trapframe */

trap:3351

- trap (struct trapframe *tf)
- sets `proc->tf` to the current trapframe on system call at 3356
- exec sets `“proc->tf->eip”` to new executable `“main”` at 6396
- exec sets `“proc->tf->esp”` to new stack at 6397

exec: 6310

`exec (char *path, char **argv)`

- read program headers from file
- load all sections
- allocate stack and copy arguments
- adjust process size
- rewrite eip and esp in trapframe
- call switchvm to load new page table

switchvm: 1873

- disable interrupts
- setup TSS to point to the process kernel stack
- load the new page table
- restore the original interrupt flags

exec: 6310

`exec (char *path, char **argv)`

- read program headers from file
- load all sections
- allocate stack and copy arguments
- rewrite eip and esp in trapframe
- adjust process size
- call `switchvm` to load new page table
- call `freevm` to free all the user pages in the old page directory

freevm: 2015

- free all user pages
- free all page table pages
- free the page directory

Virtual address space after exec



allocate space for heap

- malloc uses sbrk
- sys_sbrk (int n) : 3701
 - growproc(int n) : 2535
 - adjust the process size by n bytes
 - if n is positive
 - call allocvm
 - if n is negative
 - call deallocvm
- Why does “growproc” call “switchvm”?

Process address space



Does OS has to map the entire RAM in the kernel address space?

- Linux reserves virtual addresses between 3GB – 4GB for kernel
- Windows reserves 2GB – 4GB for the kernel (similar to xv6) but can be configured to use 3GB – 4GB

Page fault

- Page fault exception is raised by the hardware when a virtual address is dereferenced without sufficient privilege or no physical address is mapped corresponding to the virtual address
- Page fault is useful for copy on write optimization
 - The OS can make a copy, reinstate the write privilege and resume the application
- The hardware restart the execution of the partially executed instruction after returning from the exception handler

Demand paging



Demand paging

- Maintain another data structure corresponding to every page table (say frame table)
- Corresponding to every virtual page store the file identifier and offset in the frame table
- On page fault allocate a physical page, read data from the given frame offset in the frame table, and map the page in the page table