

Midterm


- You have to bring a hard copy of xv6-code-listing to the examination hall
- You can bring any additional reference (books, notes, etc.)
- Electronic gadgets are not allowed

Boot loader

- When the PC starts, it loads the boot loader into a predefined location in physical memory
- The boot loader loads the kernel text and data at the physical address 0x100000 (1 MB)
- The kernel code is compiled to be loaded at location 0x80100000
 - Why does boot loader not load the kernel at 0x80100000?

Boot loader

x v c





entry:1144

- After loading the kernel the boot loader jumps to entry:1144
- Notice that entry:1144 is the kernel code and it is compiled for addresses greater than 2 GB
- The entry:1144 loads a page table that maps kernel at 0x80100000
 - How do we load a page table when the current EIP is a physical address?


entry:1144

entry:1144

- entrypmdir:1411



$$EIP = (1MB - 4MB)$$



0x100000 → mov .i.eax, .i.cr3
0x100000 ✓
0x80100000 ✓

entrypgdir:1411

- Maps 4 MB physical page (0 – 4MB), at virtual addresses 0 and 0x80000000 (2GB)

entry:1144

1167: char stack[STACKSIZE]
1158: *esp = &stack[STACKSIZE]

- entry loads entrypgdir in the cr3 register
 - At this point the paging hardware is active
 - The virtual pages in the range 0:0x400000 and 0x80000000:0x80400000 refer to the same physical pages
 - The entire xv6 kernel and data can fit into 4 MB
 - Let us call 0:0x400000 range as low addresses
 - Let us call 0x80000000:0x80400000 range as high addresses
- Even after loading the page table the current EIP is still a low address


entry:1144

- entry uses global variable “**stack**” as a kernel stack
- entry loads the address (**stack+KSTACKSIZE**), which is a high address, in register **%esp**
- entry loads the address of the **main** in **%eax** register, which is a high address
- entry does an indirect jump through **%eax** register to jump to **main**
- After this step all the addresses are high address

Creating the first process



- userinit : 2502
- allocproc: 2455
 - allocate a process control block (PCB)
 - size, kernel stack, page directory, context, pid, trap frame, state, open files, etc.
 - allocate kernel stack
 - allocate space on kernel stack for trap frame and process context

kernel stack in allocproc



userinit : 2502

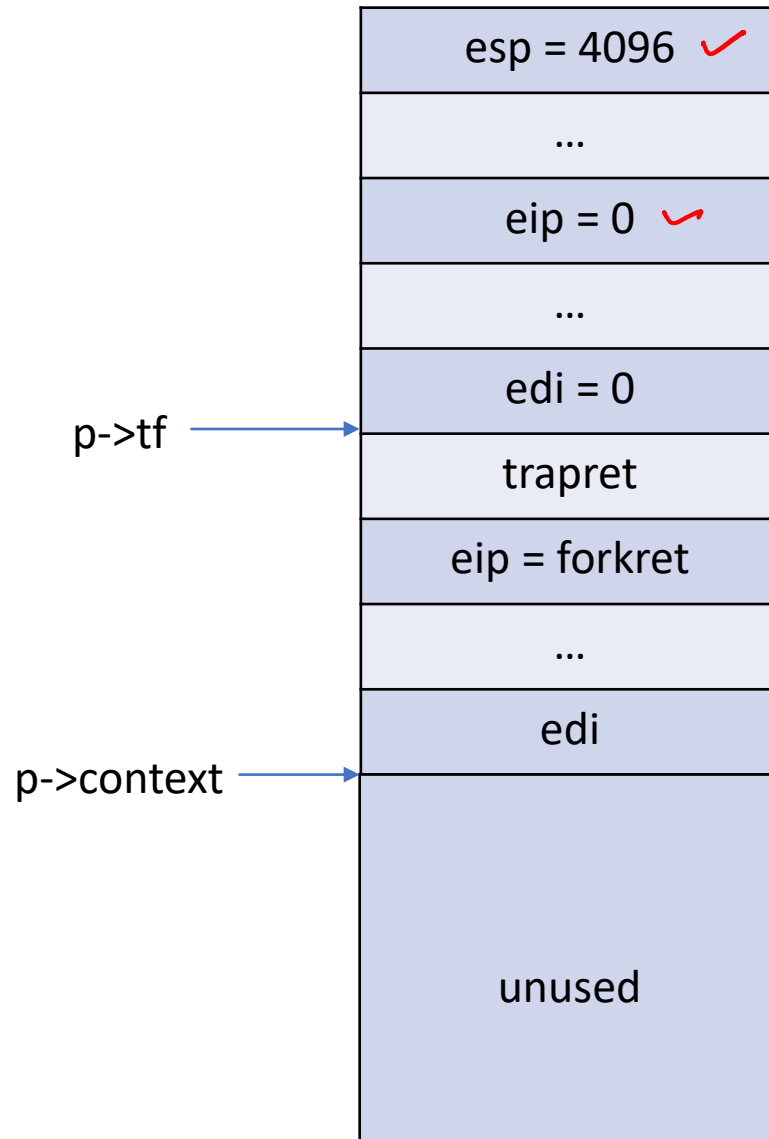
- create page directory and map kernel pages
- load “**initcode.S**” at the first virtual page in the user address space
- set up the trapframe



initcode.S:8408

- Doesn't take any argument
- Does exec system call to load “**init**” executable
- Entire code, data, and stack of **initcode.S** fit into 4096 bytes

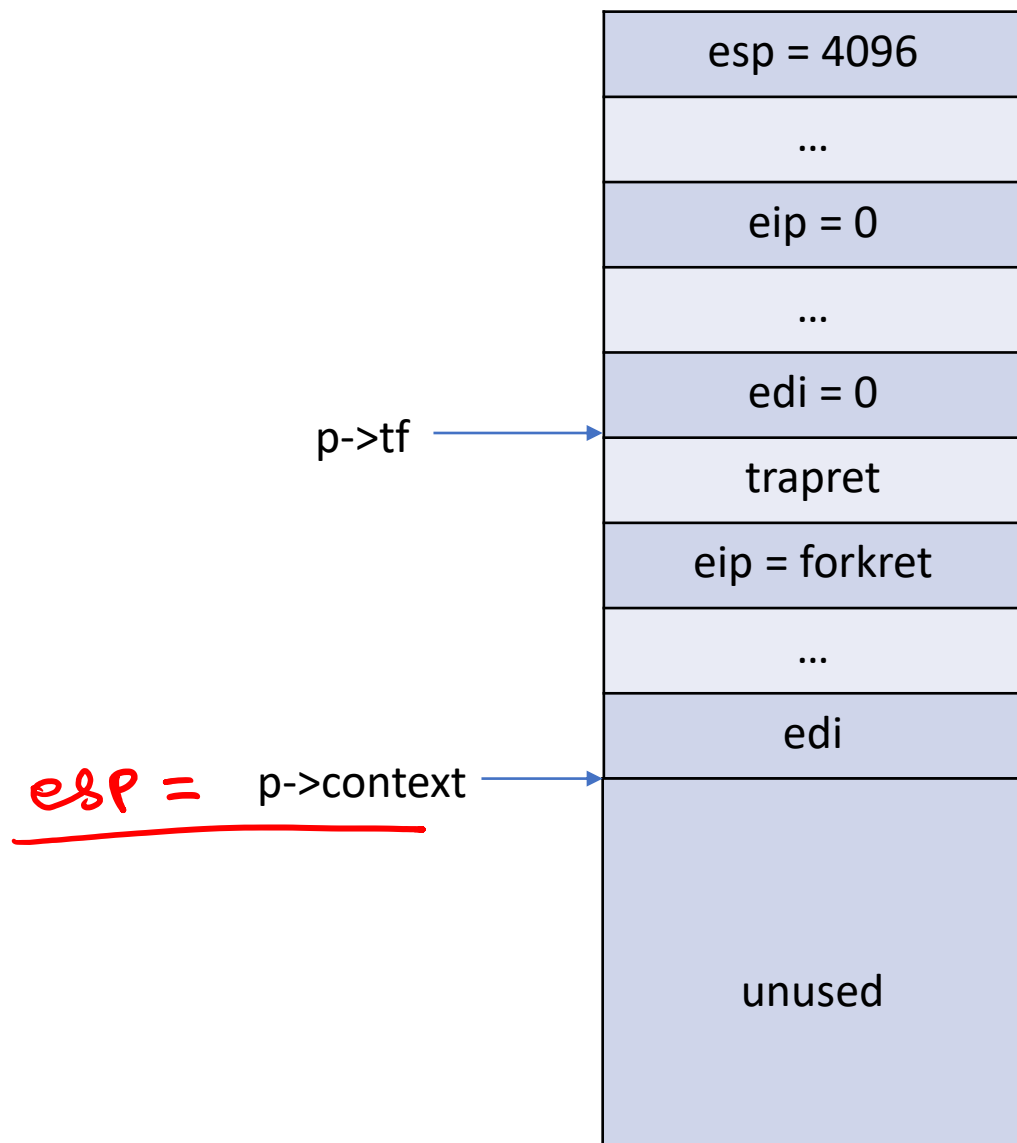
kernel stack after userinit



Schedule

- After `userinit` at some point `main` calls `mpmain` that eventually calls the `scheduler:2708`
- `scheduler:2708` finds a runnable process, switch the page table and call `swtch:2958` to load the new process context

switch:2958



switch (scheduler, target process
struct context {

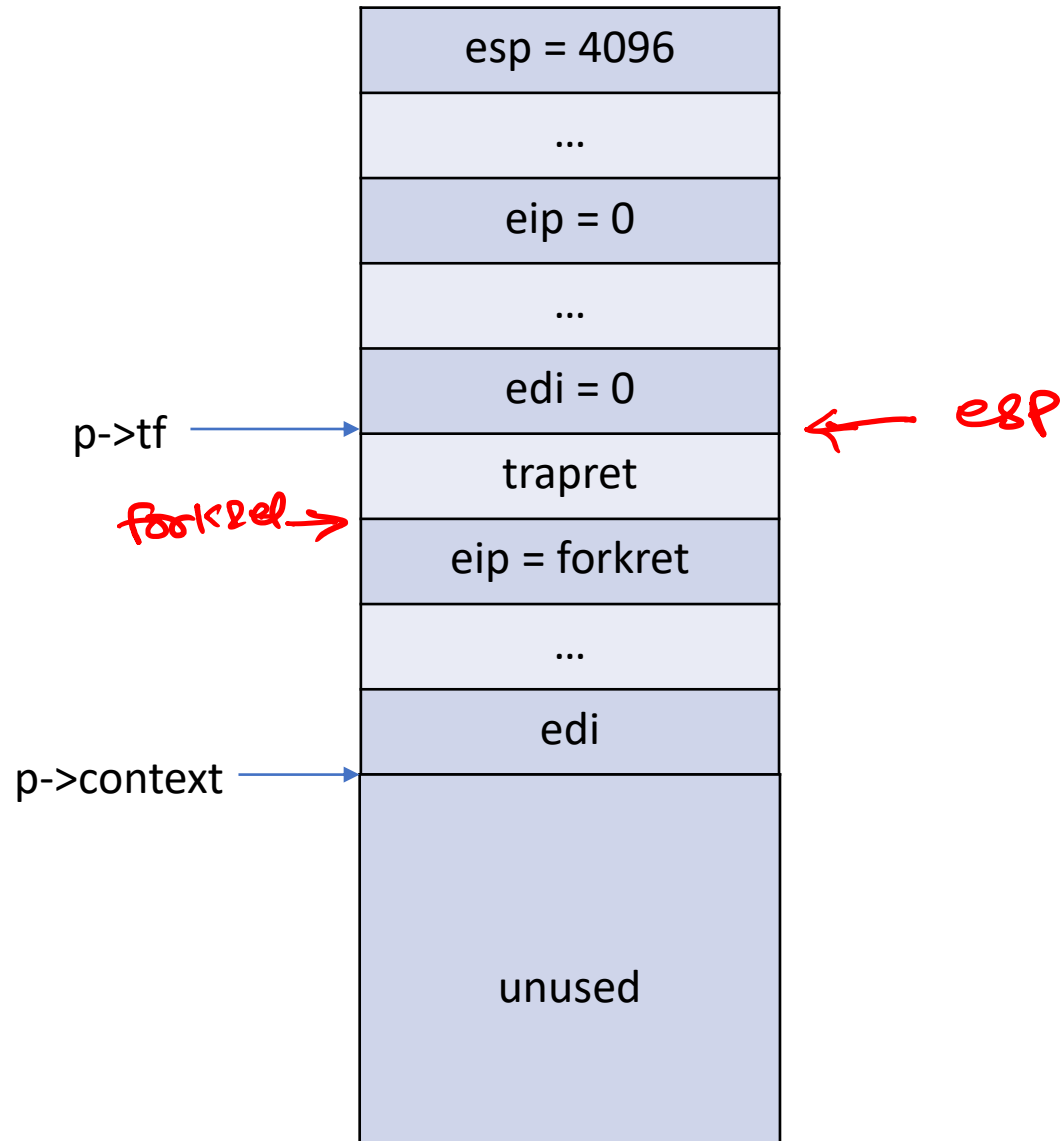
edi ;
esi ;
ebx ;
ebp ;
eip ;
}

edi, esi, ebx, ebp

forkret:2788

- Release scheduler lock and return

forkret



trapret:3277

- Pops the trap frame and executes iret

Schedule

- How does `scheduler:2708` regain control after calling `swtch`
 - call to `swtch` from `scheduler` saves the scheduler context in `cpu->scheduler`
 - `yield:2777` calls `sched:2758` that saves the current process context in “`proc->context`” and loads the scheduler context from `cpu->scheduler`