



# Locking

- Two operations
  - acquire
  - release
- The code between acquire and release is called critical section
- Locking ensures mutual exclusion
  - i.e., two threads cannot execute the critical section at the same time

# Locking

acquire ( )

Critical Section

release ( )

# Locking

- We have discussed uniprocessor locking so far
- We will discuss multiprocessor locking in future lectures
- For now, you can assume the implementation of the acquire and release as black box
- You can directly use these interfaces to ensure mutual exclusion

# Lock interfaces with lock variables

- `sema_down`
- `spin_lock`
- etc.

# Lock interfaces with lock variables

```
struct lock a, b;
```

## **thread 1:**

```
lock_init (&a, 1);  
lock_acquire (&a);  
critical_section1();  
lock_release (&a);
```

## **thread 2:**

```
lock_init (&b, 1);  
lock_acquire (&b);  
critical_section2();  
lock_release (&b);
```

**critical\_section1()** and **critical\_section2()** can execute in parallel.

# Scheduling in xv6

- Sleep
  - waiting for an event
  - waiting for I/O
  - sleep system call
- Process has completed its time slice
  - e.g., on a timer interrupt


# Challenges

- When to do context\_switch
  - timer interrupt
  - sleep
- How to support concurrent calls to scheduler
  - xv6 uses locks



# context\_switch

Figure 5-1.



# Scheduler


- Scheduler maintains a list of all processes
- In xv6, all processes are single threaded
- Scheduler maintains a list of process control blocks (PCB)
- PCB contains a pointer to the process context

# Process context

- Each thread has its own set of CPU registers, stack, and instruction pointer
- But the processor has only one set of CPU registers and instruction pointers
- The registers are multiplex among all the threads

# Process context

- The registers are multiplexed by saving them in the process context
- The EIP is also saved in the process context
- On context switch, the process allocates space for the current process's context, save registers in process context, and save the address on process context in the PCB



# What values need to be saved in the process context

eax  
ecx  
edx  
esp  
ebp  
esi  
~~edi~~  
ebx  
eip

# What values need to be saved in the process context

eip

ebp

ebx

esi

edi

ecx

edx

eax

esp

swtch:2958

eip  
ebp  
ebx  
esi  
edi  
~~ecx~~  
~~edx~~  
~~eax~~  
esp

← process\_context == esp

§ PCB → context

(struct context \*old,  
struct context \*new)

§ PCB\_OLD → context,  
PCB\_NEW → context;

(struct context \*old, struct context \*new)  
{

push ebp  
push ebx  
push esi  
push edi  
\*old = esp  
esp = new

pop edi  
pop esi  
pop ebx  
pop ebp

# scheduler:2708

- scheduler is called on scheduler stack
  - scheduler is called when the first process is created

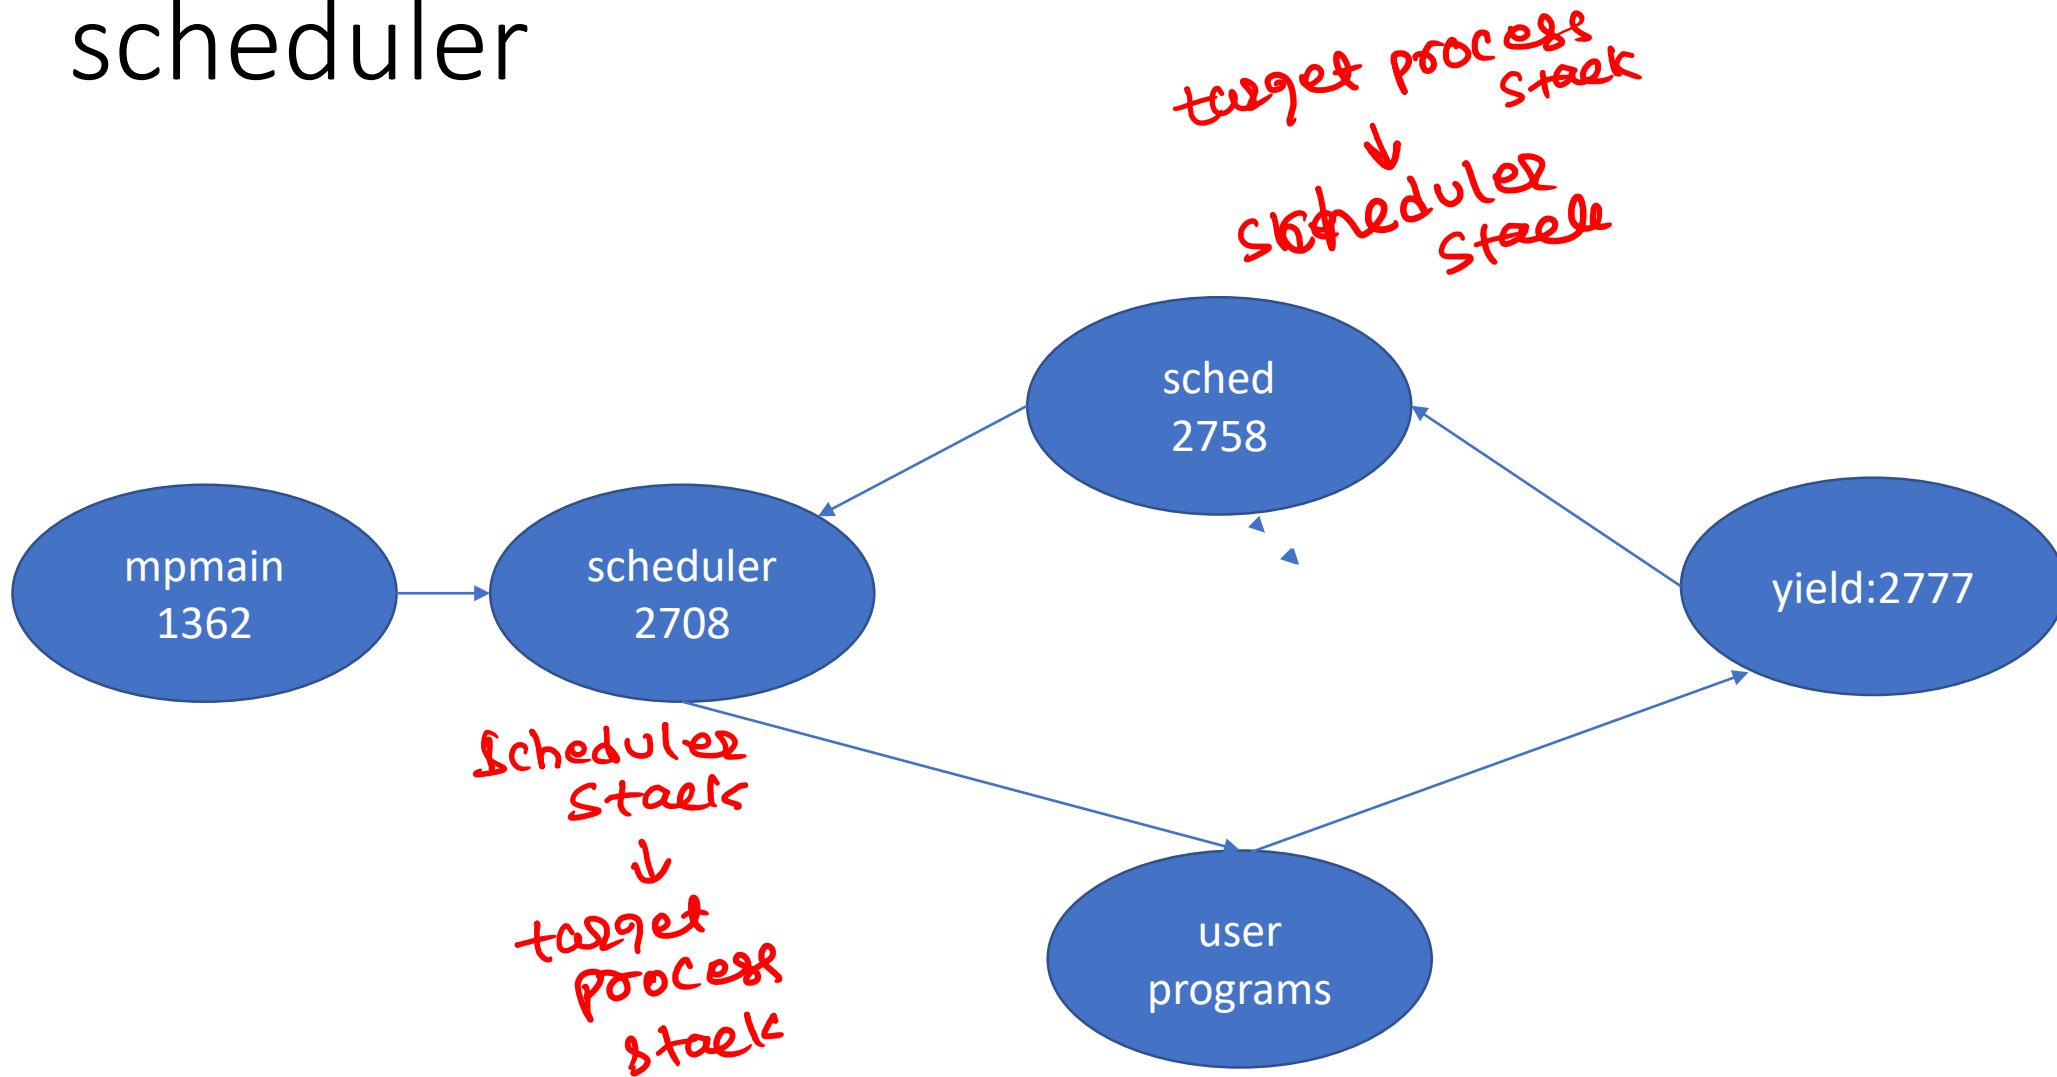




# yield:2777

- yield:2777 is called when the current process time slice expires
  - yield calls sched:2758 to save the context of current process and restore the context of scheduler

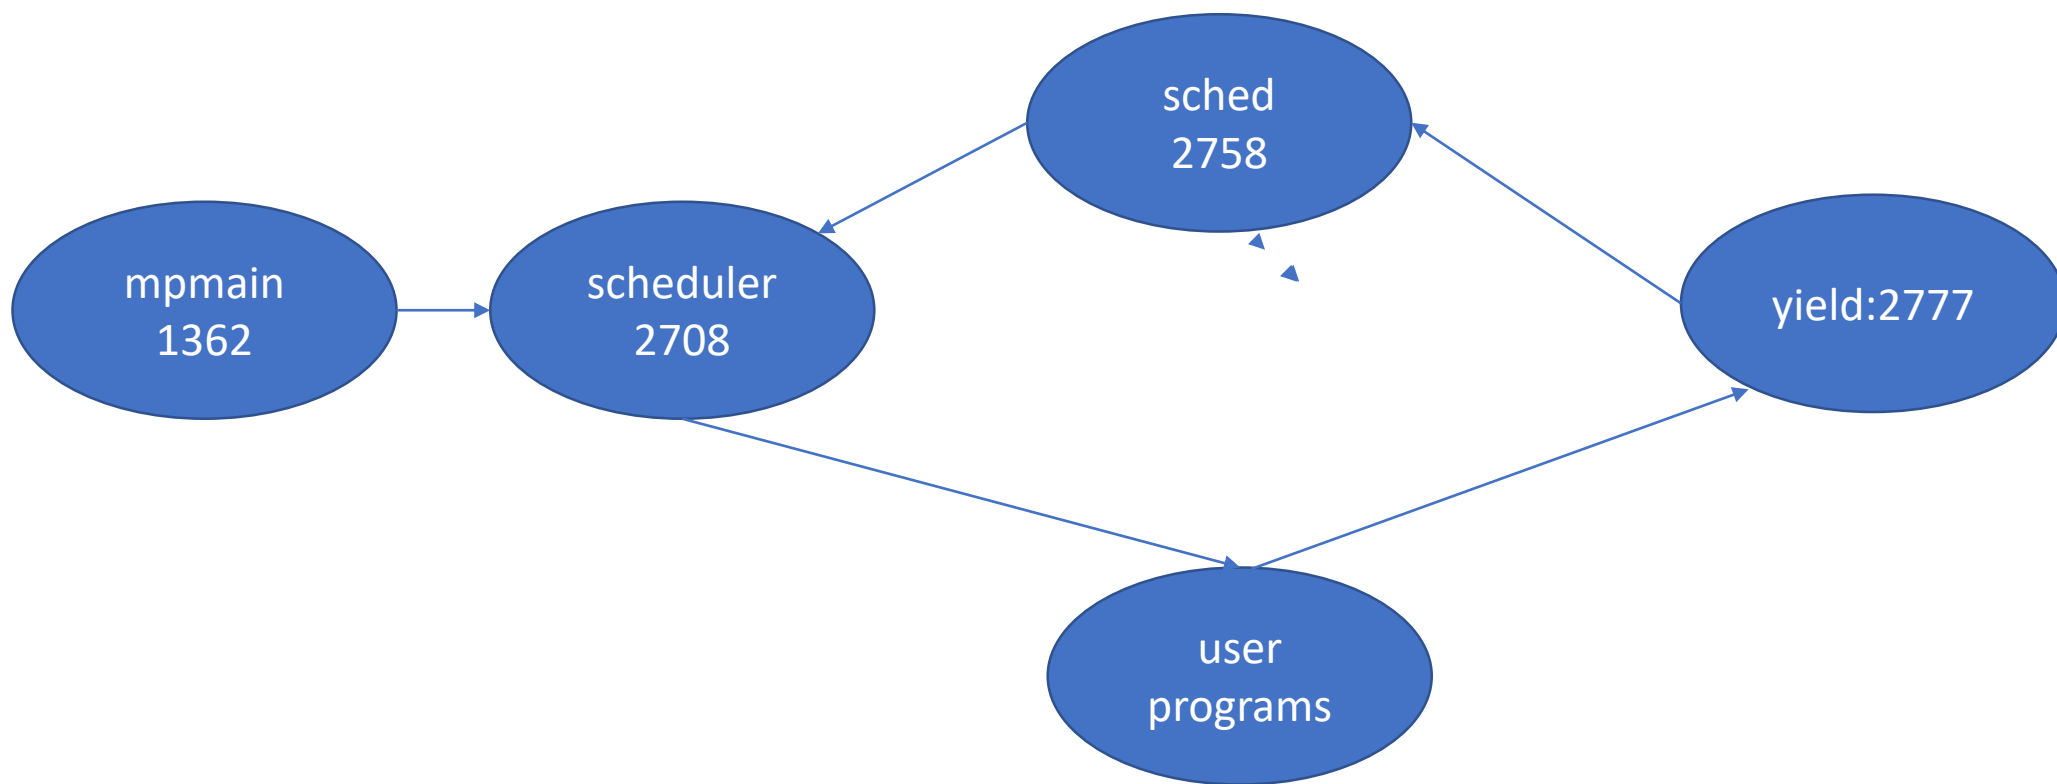
# scheduler



# PCB list

- Scheduler maintains a list of PCBs of all processes
  - Why process not thread?
- Scheduler walks to PCB list to find a runnable process for scheduling
- fork:2556 creates a new process and add it to PCB list
  - fork system call is the only way to create a new process after the first user process is loaded in the user space

# scheduler



# ptable.lock

- ptable.lock is used to protect the PCB
- what could go wrong if ptable.lock is not acquired in scheduler and yield
  - concurrent calls to yield and scheduler may try to schedule and yield the same process concurrently

```
process 1 yield
acquire (&ptable.lock);
sched(); -
return
switch (scheduler);
switch kvm();
switch (process 1);
2772
release (&ptable.lock);
```

# Concurrency without ptable lock

- “process1” does the yield system call
- yield:2780 sets the state as RUNNABLE
- scheduler:2708 is called on the other core
- scheduler finds “process1” RUNNABLE
- Both scheduler and sched call swtch to restore/save “process1” context at the same time

# What is cpu->scheduler

- CPU1:

yield()

acquire (&ptable.lock)

sched:

load cpu->scheduler context

schedule process<sub>1</sub>

save cpu->scheduler context

release (&ptable.lock)

- CPU2:

yield()

acquire (&ptable.lock);

sched:

load cpu->scheduler context

but cpu->scheduler contains  
context of CPU1

“cpu” is per-CPU variable

- “cpu” points to different memory locations for different CPUs
- mpmain() is called for each CPU that eventually calls the scheduler on each CPU



# scheduler and yield disable interrupts

- `acquire(&ptable.lock)` internally disables interrupt
  - why we need to disable interrupts in scheduler

# scheduler:2708

- Why acquire() and release() inside the outer for loop?
  - Why not outside the outer for loop?

# Sleep and wakeup

```
struct q {  
    void *ptr;  
};  
  
void *send (struct q *q, void *p) {  
    while (q->ptr != 0){  
    }  
    q->ptr = p;  
}
```

```
void *recv (struct q *q) {  
    void *p;  
    while ((p = q->ptr) == 0) {  
    }  
    q->ptr = 0;  
    return p;  
}
```

Not good if sender sends rarely!

# Sleep and wakeup

```
struct q {  
    void *ptr;  
};
```

```
void *send (struct q *q, void *p) {  
    while (q->ptr != 0) {  
    }  
    q->ptr = p;  
    wakeup (q);  
}
```

```
void *recv (struct q *q) {  
    void *p;  
    while ((p = q->ptr) == 0) {  
        sleep (q);  
    }  
    q->ptr = 0;  
    return p;  
}
```

# Sleep and wakeup

```
struct q {  
    struct spinlock lock;  
    void *ptr;  
};  
  
void *send (struct q *q, void *p) {  
    acquire (&q->lock);  
    while (q->ptr != 0) {  
    }  
    q->ptr = p;  
    wakeup (q);  
    release (&q->lock);  
}
```

```
void *recv (struct q *q) {  
    void *p;  
    acquire (&q->lock);  
    while ((p = q->ptr) == 0) {  
        sleep (q);  
    }  
    q->ptr = 0;  
    release (&q->lock);  
    return p;  
}
```

# Sleep and wakeup

```
struct q {  
    struct spinlock lock;  
    void *ptr;  
};  
  
void *send (struct q *q, void *p) {  
    acquire (&q->lock);  
    while (q->ptr != 0) {  
    }  
    q->ptr = p;  
    wakeup (q);  
    release (&q->lock);  
}
```

```
void *recv (struct q *q) {  
    void *p;  
    acquire (&q->lock);  
    while ((p = q->ptr) == 0) {  
        sleep (q, &q->lock);  
    }  
    q->ptr = 0;  
    release (&q->lock);  
    return p;  
}
```

sleep:2809

- How does sleep ensure that wakeup will not be lost?
- What is the check at 2823?

# wakeup:2864

- What could go wrong if wakeup does not acquire ptable.lock?
- Can we do it more efficiently?



# pipe

- pipe is a circular buffer of size PIPESIZE
  - pipe\_is\_empty =  $nread == nwrite$
  - pipe\_is\_full =  $nwrite - nread == PIPESIZE$

nread = number of bytes read



nwrite = number of bytes written

# piperead:6551

- piperead (struct pipe \*p, char \*addr, int n);
  - tries to read n bytes from the pipe (p) to buffer (addr)
  - returns the number of bytes read
  - if the pipe is empty wait until some bytes are written to the pipe

# pipewrite:6530

- pipewrite (struct pipe \*p, char \*addr, int n);
  - writes n bytes from the buffer (addr) to pipe (p)
  - if the pipe is full waits until a reader consumes some bytes

reader  
is  
sleeping

# exit and wait

## **exit**

- Mark the status of the current process as zombie
- If the parent has already exited, change the parent to init process
- wakeup the parent process
- call schedule

## **wait**

- If any of the children is in zombie state, releases the resources of the child and returns the child pid
- If none of the children has exited, sleep until the child process wakes up

# exit:2604

- Can exit free all the resources in this routine?
- Can wait:2653 free the child resources just after the exit:2604 sets the status of the exiting process to zombie?

# kill:2875

- kill system call allows a process to kill another process
- killing a process at arbitrary points is tricky because the target process might be in the middle of something
  - e.g., holding a kernel lock
- kill:2875 simply sets a field “killed” in the target process PCB
- trap:3351 checks the killed flag before returning to user mode
  - if the current process was killed, trap calls the exit

