

LRU implementation

- Replace a page which has not been used for the longest period
- How to identify a page that was not accessed recently?

LRU implementation

- We need some hardware support to update the timestamp on every access
- x86 provide some help in the form of a reference bit (access bit)
- The access bit is set when a page is accessed by an instruction

Access bit


- mov \$100, 0x10101 // write 100 to address 0x10101

0x10 → pre
access bit

Additional reference bit algorithm

- keeps n-bit timestamp for each page in memory
- At regular interval right shift the timestamp by one bit
- Shift the access bit to the high order bit of timestamp
- Replace all the pages with the lowest timestamp
 - or replace the next page in FIFO order

Additional reference bit algorithm



Second chance algorithm


- A simplification of the previous algorithm is second chance algorithm when we don't want to have additional storage to store the timestamp
 - i.e. n is zero
- The algorithm is similar to FIFO but gives more priority to a page which is accessed recently

FIFO

- FIFO algorithm can be implemented using a circular queue
- The size of the circular queue is the number of available physical pages
- When the queue is full, we need to select a victim that is going to be replaced
- A pointer (also called clock tick) points to the next victim in the queue
- After replacing the victim, the clock tick points to the next element in the queue


FIFO

- 7 0 1 2 3 4 5 9 10 11 12



FIFO

- 7 0 1 2 3 4 5 9 10




Second change algorithm

```
select_victim (struct queue_elem *clock) {  
    while (1) {  
        page = clock->page;  
        if (!was_accessed (page)) /* check for access bit */  
            return page;  
        clear_access_bit (page); /* reset access bit */  
        clock = clock->next;  
    }  
}
```


Second chance algorithm

- If a page was referenced recently, clear the access bit and move to the next page



Second chance algorithm

- If a page was referenced recently, clear the access bit and move to the next page



Dirty bit

- `mov $100, 0x10101` `// write 100 to address 0x10101`

0x10
`pte = pte (0x10)`
`pte->dirty = 1;`

Enhanced second chance algorithm

- In the case of memory mapped files, a page need not need to be written to the disk if it is not modified

*char *addr = mmap(15 * 1024);*

- Page table entry also contains a dirty bit
 - dirty bit is set when an instruction writes to a page

**(addr + 100)
mmap(addr);*

- The enhanced second chance algorithm gives preference to the dirty page over clean page

Enhanced second chance algorithm


```
select_victim (struct queue_elem *clock) {  
    while (1) {  
        page = clock->page;  
        if (!was_accessed (page)) { /* check for access bit */  
            if (!was_written (page) || page->given_advantage) /* check for dirty bit */  
                return page;  
            page->given_advantage = true;  
        } else {  
            clear_access_bit (page);  
            page->given_advantage = false;  
        }  
        clock = clock->next;  
    }  
}
```

~~access dirty~~
~~0 0~~

dirty	access
0	0
1	0
0	1
11	1

Enhanced second chance algorithm

- If a page was referenced recently, clear the access bit and move to the next page
- Otherwise, if a page is dirty, give it another chance




Allocation of frames

- An OS must decide how page frames are allocated for a process
- Two allocation policies
 - Local
 - Global

Local allocation

- In the local allocation scheme, each process is allocated a fixed set of physical frames
- The process must allocate physical pages from its assigned quota of physical frames
- One problem of this approach is underutilization of memory
- Number of page faults are deterministic for a given workload


Local allocation



Global allocation


- OS allows processes to allocate from the set of all physical frames
- One process can cause replacement for other processes
- Number of page faults are not deterministic
- Good throughput (no underutilization)

Global allocation




Working set

- The working set of a process is the set of pages accessed in last Δ page references
- For example if $\Delta = 10$



Thrashing


- If the sum of working sets of active applications is larger than the total number of physical frames, then the applications will spend most of the time in swapping
- This results in very low CPU utilization because most of the time will be spent in disk I/O
- This is also called thrashing



Thrashing

- To avoid thrashing the scheduler maintain two lists
 - active list
 - inactive list
- Applications may move between active and inactive list
- The scheduler picks a process from the active list
- Processes are added to the active list in such a way that the sum of working sets of active processes is less than the total number of physical frames

Thrashing



xv6

```
exec {  
    va = kalloc(  
    read(va, executable, size);
```

Other uses of page protection


- Temporal safety

Temporal safety

```
struct List {  
    int data;  
    struct List *next;  
};
```

```
struct Point {  
    int x;  
    int y;  
};
```

```
delete_node (struct List *l, int val) {  
    while (l) {  
        if (l->data == val)  
            free (l);  
        l = l->next;  
    }  
}  
  
add_coordinate (int x, int y) {  
    struct Point *p = malloc (sizeof(*p));  
    p->x = x; p->y = y;  
    return p;  
}
```



Temporal safety

```
struct List {  
    int data;  
    struct List *next;  
};
```

```
struct Point {  
    int x;  
    int y;  
};
```

```
delete_node (struct List *l, int val) {  
    while (l) {  
        /* may access a dangling reference */  
        if (l->data == val)  
            free (l);  
    }  
}  
  
add_coordinate (int x, int y) {  
    struct Point *p = malloc (sizeof(*p));  
    p->x = x; p->y = y;  
    return p;  
}
```

Temporal safety

- Never reuse a virtual address
 - malloc always returns a new virtual address
- Free does nothing
- Why does this scheme work?
- What is this scheme not efficient?

On free take back physical page?

- Does this scheme work?

Allocate virtual pages for every allocation

VP1 = malloc (20); // VP1 is a virtual page of 4096 bytes

VP2 = malloc (30); // VP2 != VP1 and VP2 is a virtual page of 4096 bytes

On free take back physical page?

- Throw dangling pointer exception if a page fault is encountered
- Does this scheme work?
 - Yes, but requires a lot more physical memory because every allocation needs to be aligned to the page size

Can we do better?

- Share physical pages

Share physical pages

- Allocate different virtual pages for every allocation but share the same physical pages

e.g.,

`VP1 = malloc (20);` `// VP1 points to PP1 (physical page)`

`next malloc (30) returns VP2 + 20.` `// VP2 also points to PP1 and VP2 != VP1`

Drawbacks?

- Fragmentation
- TLB pressure