

Double free

```
char *ptr = malloc (100);
```

```
...
```

```
free (ptr);
```

```
...
```

```
free (ptr);
```

```

void* mymalloc (int size) {
    char *ptr = malloc (size + 4);
    if (ptr != NULL) {
        ptr[0] = 1;
        ptr += 4;
    }
    return (void*)ptr;
}

```

```

void* myfree (void *_ptr) {
    if (_ptr != NULL) {
        char *ptr = (char*)_ptr - 4;
        if (ptr[0] != 1) {
            printf ("double free detected\n");
            exit (0);
        }
        ptr[0] = 0;
        free (ptr);
    }
}

```

- The application code is using **mymalloc** and **myfree** APIs
- **malloc** and **free** are implementing buddy allocation
- Will **myfree** always terminate a program if the program has double free bug?
- Justify your answer

Hash function

- Maps data of arbitrary size onto a data of fix size

```
char buf[] = "Hello world\n";
```

```
-----
```

```
int hash = 0;
```

```
int i;
```

```
for (i = 0; i < strlen (buf); i++)
```

```
    hash += (int)buf[i];
```

```
return hash;
```

Hash function

- A simple hash function has collisions
- If we use the hash function of the previous slide “Hello world\n” and “world Hello\n” yield the same hash value

Cryptographic hash function

- It is infeasible to find two messages whose hash values (generated by cryptographic hash function) are same
 - e.g., “Hello world\n” and “world hello\n” will have very different hash values
- **SHA1** generates 20 bytes hash value for any arbitrary length of string
 - It is infeasible to find two messages whose **SHA1** are same

Integrity of files

- Suppose a process wants to protect the **integrity** of files created by it
- By integrity, we mean that no other process is allowed to modify the file contents

Integrity

- Assuming that a process has a special file (say `secure.txt`) that can not be tampered
- Store SHA1 of the file contents in `secure.txt`

read

- read the entire file
- compute the SHA1 of the file contents
- check the SHA1 against the one stored in the **secure.txt** file
- report an error if the SHA1 does not match

write

- read the entire file
- check the integrity
- update the data in the read buffer
- recompute SHA1
- update SHA1 in **secure.txt**
- update the file

Problems with the previous approach

- On every *read/write the entire file on every read/write* read/write, we need to read the entire file

Alternative approach

- Divide the files into blocks
- Compute SHA1 at the block granularity
- Store the SHA1 for each block in **secure.txt**
- On every read/write, read the conflicting blocks and check the integrity

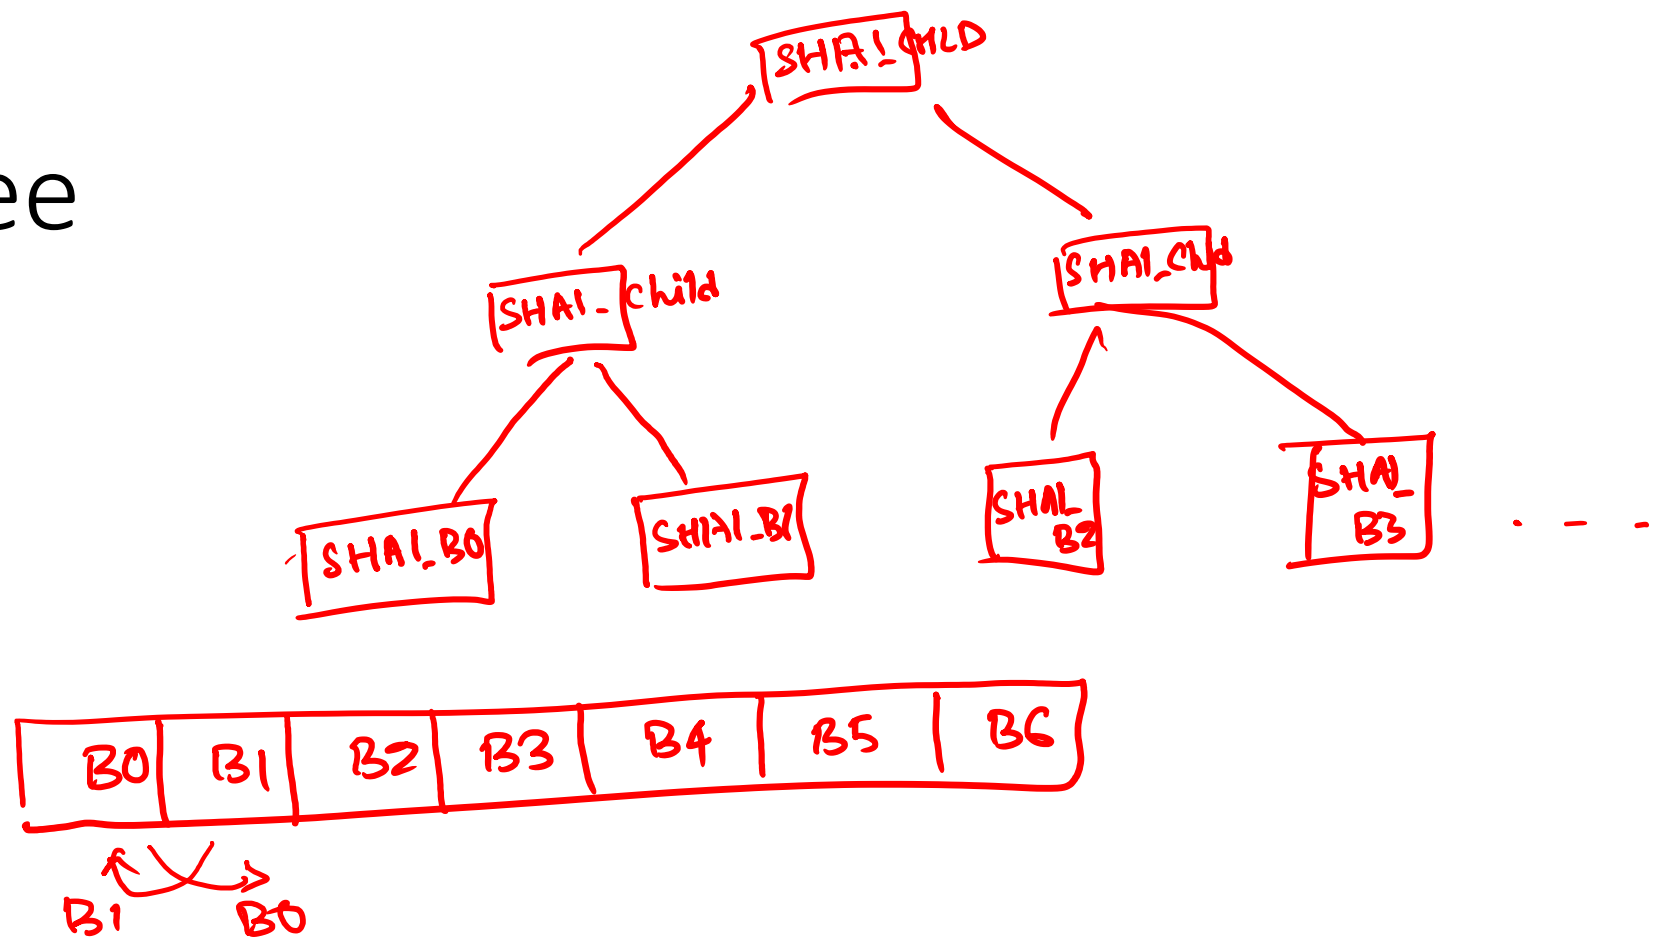
Problem with this approach

- Number of entries in the `secure.txt` depends on the size of the file

Merkle tree

- File is divided into blocks
- The leaf nodes of Merkle tree contain the **SHA1** of the blocks
- An internal node in Merkle tree contain the **SHA1** of the concatenation of its child nodes
- The root of the Merkel tree is unique for a given file

Merkle tree



Merkle tree

- What if we reorder two blocks?

concatenation will ensure the integrity

Next assignment

- Protect the integrity of files created by a library that implements Merkle tree