

Cache

0x1000
↓
mov x, .1.eax

Instructions

Cache

0x1000-

Main memory

Cache

- Cache is a fast storage media between RAM and CPU
 - Cache is ten times faster than RAM
- data is brought to cache before it can be read/written by the instructions
- typical cache line size is 64 bytes
- an entire cache line is brought to the cache when a memory byte is accessed

Write-back cache



$x = 20$

$x = 20$

Write-back cache

mov \$10, x

x = 10

x = 20

When a cache line is written to main memory?

- Asynchronous
 - At the time of replacement

How a memory cache write-back is different from buffer cache write-back?

buffer cache

- write-back in buffer cache does not necessarily mean replacement
 - write-back happen during commit
- The commit policy also depends upon how an application interacts with the user

Write-through cache



$x = 20$

$x = 20$

Write-through cache

mov \$10, x

x = 10

x = 10

Disadvantage of write-through

- Every write goes to RAM
 - slow

Advantages of write-through


Advantages of write-through

- Cache coherence hardware is simple


Can directly read write with memory

- Memory mapped devices require write-through memory
 - Page table entry contains a write-through flag

Write-back cache



Write-back cache



Write-back cache

CPU1

mov x, %eax

x = 20

CPU2

mov \$10, x

x = 10

x = 20

Cache coherence

- A cache coherence protocol ensures the consistency of shared data among multiple per-processor caches

MESI

- A cache line can be in any of these states
- **Modified (M)**: doesn't match main memory and only present in current cache
- **Exclusive (E)**: matches main memory and only present in current cache
- **Shared (S)**: matches main memory and may present in multiple private caches
- **Invalid (I)**: this cache line is invalid

MESI

- For a pair of caches, the permitted states of a give cache line as follows:

	M	E	S	I
M	x	x	x	y
E	x	x	x	y
S	x	x	y	y
I	y	y	y	y

MESI

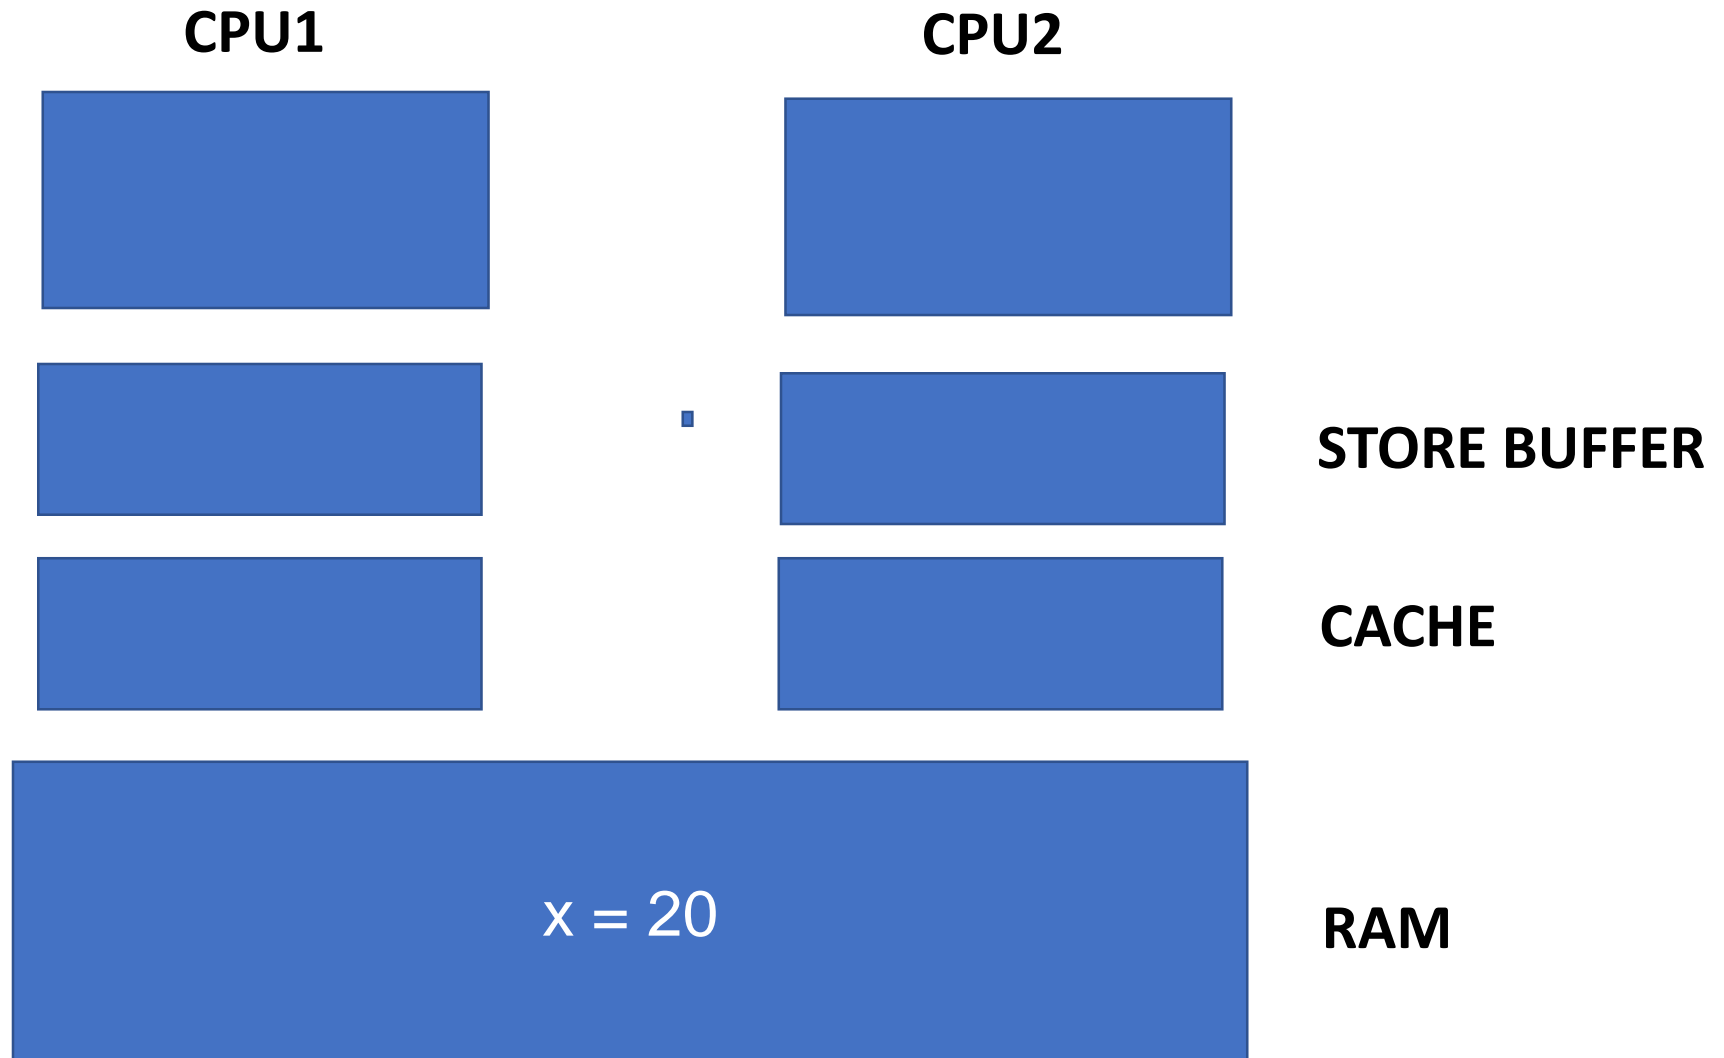
M - Modified
E - Exclusive
S - Shared
I - Invalid

REQ	P1	P2	P3
Initially	NP	NP	NP
R1	E	NP	NP
W1	M	NP	NP
R3	S	NP	S
W3	I	NP	M
R1	S	NP	S
R3	S	NP	S
R2	S	S	S

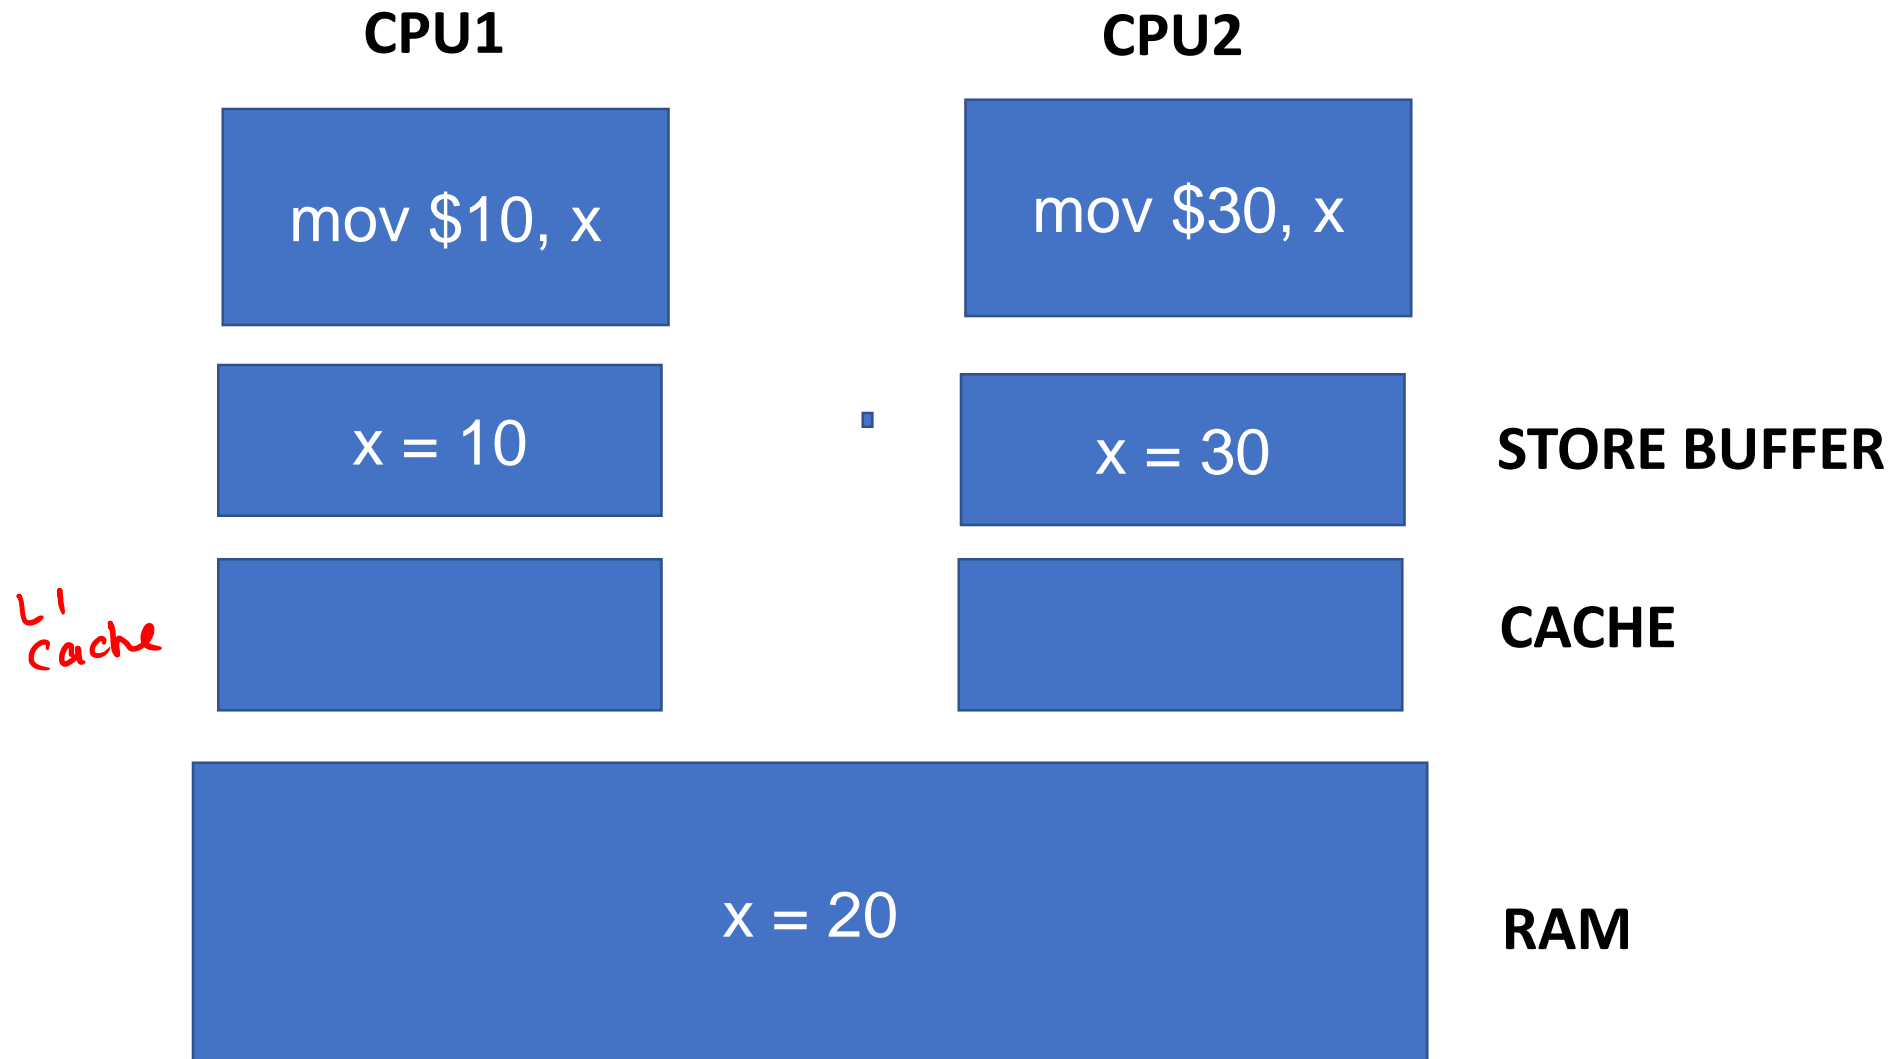
write back

write back

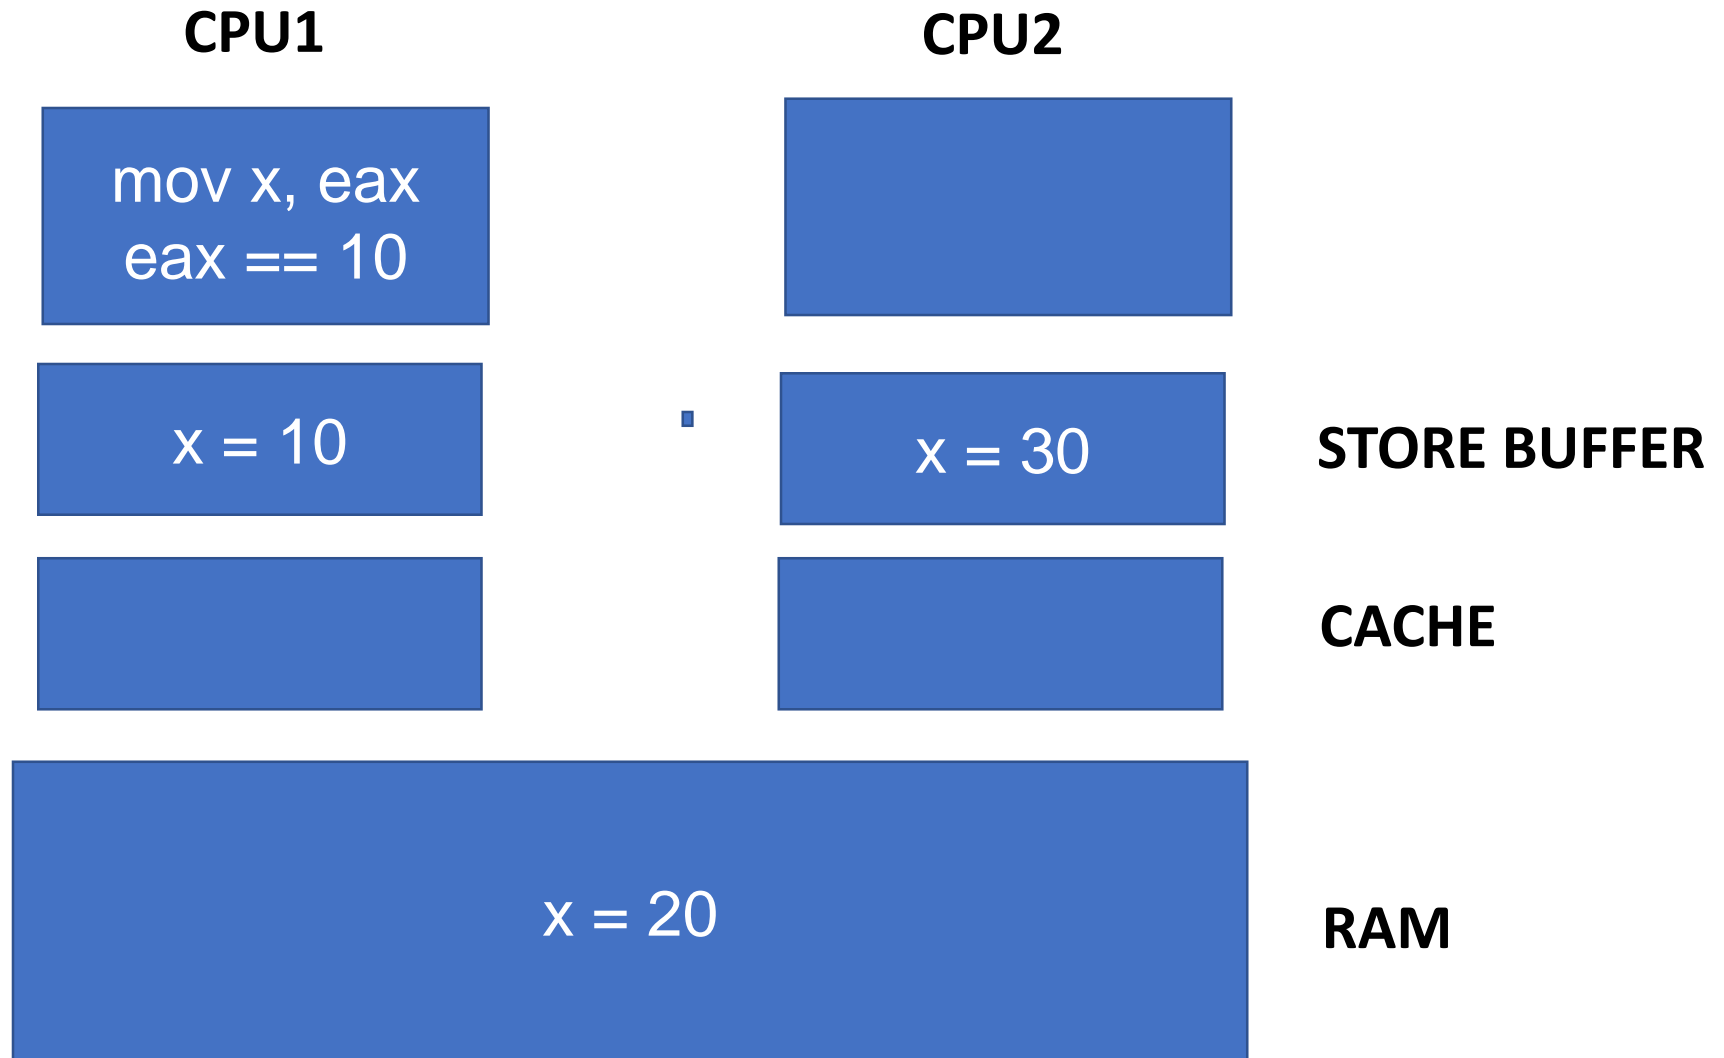
Store buffer



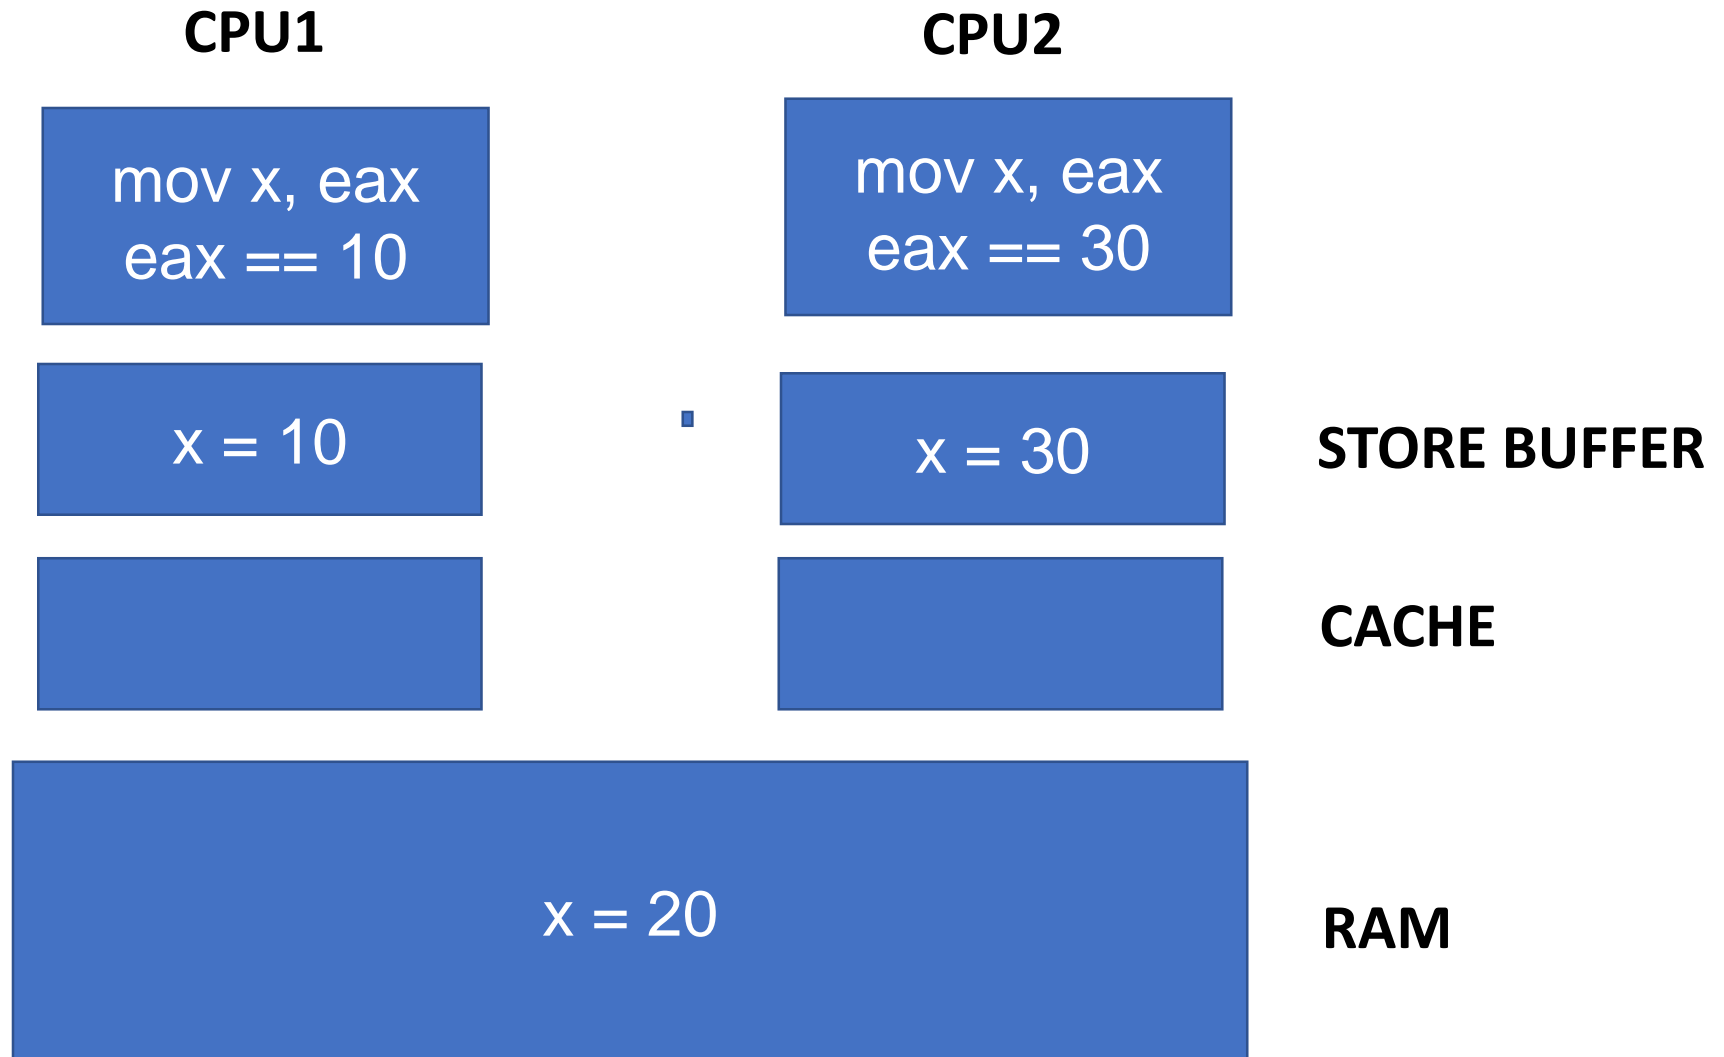
Store buffer



Store buffer



Store buffer



Store buffer

- No cache coherency at the store buffer granularity
- reads from the same core, which are present in the store buffer, are served from the store buffer

Memory ordering

- load and store are atomic operations on x86
- Memory ordering refers to the order in which the processor issues read and write request to the system memory

$x = 1$ $x = 2$
 $x = 1 \mid 2$

Memory ordering

P0	P1
✓ x = 1	y = 1 ✓
✓ <u>r1</u> = y	r2 = x ✓

Initially, x and y are 0.

Handwritten notes illustrating memory ordering scenarios:

Scenario 1 (Top Left):

```
x = 1
r1 = y
y = 1
r2 = x
```


Scenario 2 (Top Right):

```
x = 1
y = 1
r1 = y
r2 = x
```

Scenario 3 (Bottom):

```
r1 = y
x = 1
y = 1
r2 = x
```

on x86, loads and stores are not reordered
with like operations

P0	P1
 <code>x = 1</code>	<code>r1 = y</code> 
<code>y = 1</code>	<code>r2 = x</code>

`r1 == 1`
`r2 == 0`

Initially, x and y are 0.

Stores are not reordered with earlier loads

P0	P1
r1 = x	r2 = y
y = 1	x = 1

r2 == 1 & r1 == 1

Initially, x and y are 0.

Loads may be reordered with earlier stores to a different memory location

P0	P1
x = 1 <i>wfence</i>	y = 1
r1 = y	r2 = x

r1 == 0 y r2 == 0

Initially, x and y are 0.

Memory reordering

- On x86, a **load** may be reordered with an **older stores** to a different memory location
- a **load** may not be reordered with a **locked** instruction

For details you can refer to Chapter-11 of **Intel software developers manual vol-3**

How to prevent reordering?

- memory barrier (mfence) on x86
- a memory barrier pause execution until all the loads and stores prior to memory barrier are completed
- the store buffer is drained on memory barrier

Locking on multiprocessor

- lock_acquire
- lock_release

Peterson's solution

volatile int turn;

volatile boolean flag[2];

acquire: $flag[j] = false$

flag[i] = TRUE;

turn = j; *inference*

while (flag[j] && turn == j);

a load can be reordered
with earlier stores
on the different loc

release:

flag[i] = FALSE;

store flag[i]
load j
store turn
load flag[i]
load turn
load j

load j
load flag[i]
store flag[i]
store turn
load j
store flag[i]
store turn
load flag[i]

Peterson's solution

- Does not work on multiprocessor because of reordering

```
volatile int turn;  
volatile boolean flag[2];
```

acquire:

```
flag[i] = TRUE;
```

```
turn = j;
```

__sync_synchronize ();

```
while (flag[j] && turn == j);
```

release:

```
flag[i] = FALSE;
```

__sync_synchronize

- is a memory barrier (**mfence** on x86)
- no reordering across memory barrier
- drains the store buffer
- __sync_synchronize waits for all the outstanding read/write operations to finish

lock prefix

~~add~~ \$t, 0x1000

$t = *0x1000$

$t = t + 1$

$*0x1000 = t$

- certain instructions can be prefixed with “lock”
- “lock” prefix ensures the atomicity of instruction
 - x86 instructions are complex, and a single instruction can do multiple reads and writes
 - you can think of “lock” prefix as: it inserts an acquire/release around the instruction
- “lock” prefix works directly on the cache bypassing store buffer
- locked instruction drains the store buffer on current CPU

lock prefix

add \$1, 0x1000

0x1000 = 0;

Thread 1:

for 1000 times acquire

add \$1, 0x1000

release

Thread: 2

for 1000 times acquire

add \$1, 0x1000

release

lock add \$1, 0x1000

lock prefix

add \$1, 0x1000

0x1000 = 0;

Thread 1:

for 1000 times

lock add \$1, 0x1000

Thread: 2

for 1000 times

lock add \$1, 0x1000