

Optional refresher module assignment

- `send()`: wait until the entire buffer is sent
- `receive()`: wait until some bytes are available
- `send()` and `receive()` recycle the shared memory (e.g., can implement circular buffer)
- The shared memory can be accessed only through the send and receive APIs

Group assignments

- All the group members are expected to understand the complete assignment
- This will also help you in preparing for the final examination

Difference between `log_write` and `bwrite`?

Difference between log_write and bwrite?


- bwrite ensures the atomic write of single disk sector
- log_write ensures the atomic write of a sequence of disk blocks
 - i.e., the atomicity of a complete operation

Does `end_op` immediately commit?

Can we commit just once, before the program termination?

Cache lines


- VA \rightarrow PA
- [4096 – 8192] \rightarrow [0 – 4096]
- Which cache line would be brought in the cache when the application access a virtual address (4096 + 8)



Does x86 reorder loads and stores?

Does x86 reorder loads and stores?

- Yes, on x86 a load can be reordered with a previous store on a different memory location



Initially, $x == 0$ $y == 0$

(~~$x == 0$~~ $r1 == 0$) $\&\&$ $r2 == 0$)

Why does x86 reorder loads and stores?

Efficiency

Do we need **mfence** in a single threaded application?

Peterson's solution

```
volatile int turn;
```

```
volatile boolean flag[2];
```

acquire:

→ flag[j]

```
flag[i] = TRUE;
```

```
turn = j;
```

inference →

```
while (flag[j] && turn == j);
```

release:

```
flag[i] = FALSE;
```

lock prefix

- Some instructions can additionally take a lock prefix to tell the hardware to execute the instruction atomically

lock add \$1, 0x1000

Because of the lock prefix, the hardware will execute this instruction atomically

What does lock prefix do?

- drains the store buffer
- lock the cache-lines used by the instruction before executing the instruction
- loads/stores cannot be reordered across lock instruction

Disadvantage of using lock prefix

- Very slow
 - implicit memory barrier
 - lock the cache-lines (slow the execution on other cores with conflicting memory access)
- Aggressively invalidate cache lines of other CPUs if all CPUs try to modify the same cache line
 - also called cache line bouncing
 - very slow if you have large number of cores (≥ 32)

Cache line bouncing

label1:

lock add \$1, lockvar


cmp \$1, lockvar

je label2

lock sub \$1, lockvar

jmp label1

label2:



On every update : invalidate cache line
On every read : fetch the cache line
in shared mode

Spin lock

```
struct spinlock {  
    volatile unsigned int locked;  
};
```

```
void acquire (struct spinlock *lk) {  
    while (lk->locked != 0);  
    lk->locked = 1;  
}
```

```
void release (struct spinlock *lk) {  
    lk->locked = 0;  
}
```

Spin lock

```
struct spinlock {  
    volatile unsigned int locked;  
};
```

```
void acquire (struct spinlock *lk) {  
    pushcli();  
    while (atomic_xchg(&lk->locked, 1) != 0);  
    __sync_synchronize ();  
}
```

*mov \$1, %eax
lockxchg %eax, %lk → lockob*

Spin lock

```
struct spinlock {  
    volatile unsigned int locked;  
};
```

```
void release (struct spinlock *lk) {  
    __sync_synchronize ();  
    lk->locked = 0;  
    popcli();  
}
```

Spin lock

- pushcli and popcli ensures that this lock primitive can be used in an interrupt handler
- It is unsafe to use a spin lock, that does not disable interrupt, in an interrupt handler

Spin lock

- Can we improve this spin lock implementation?

```
void acquire (struct spinlock *lk) {  
    pushcli();  
    while (atomic_xchg(&lk->locked, 1) != 0) {  
        __sync_synchronize ();  
    }  
}
```


while (lk->locked != 0);

Spin lock

- What is the problem with this spin lock?
 - not fair, anybody can acquire the lock irrespective of their arrival

```
void acquire (struct spinlock *lk) {  
    pushcli();  
    while (atomic_xchg(&lk->locked, 1) != 0) {  
        while (lk->locked == 1);  
    }  
    __sync_synchronize ();  
}
```

ticket spin lock



- **atomic_xadd**: atomically adds the input value to the input memory location and return the old value of the memory location

ticket spin lock

```
struct lock {  
    volatile unsigned head; // initially 0  
    volatile unsigned tail; // initially 0  
};
```

acquire:

```
oldtail = atomic_xadd (&lockvar->tail, 1);  
while (oldtail != lockvar->head);
```

release:

```
lockvar->head++;
```

head == 0
tail == 0

Thread 1

tail == 1
oldtail == 0

Thread 2

tail = 2
oldtail = 1
while (oldtail != head)

Thread 3

tail = 3
oldtail = 2
while (oldtail != head);

head == 1

Readers-writer locks

- Multiple readers can concurrently execute in the critical section
- Only a single writer is allowed in the critical section

Readers-writer locks

```
volatile unsigned lockvar = 10000;
```

read_acquire:

```
while (atomic_sub(&lockvar, 1) < 0) {  
    atomic_add (&lockvar, 1);  
    while (lockvar <= 0);  
}
```

read_release:

```
atomic_add (&lockvar, 1);
```

write_acquire:

```
while (atomic_sub (&lockvar, 10000) != 0) {  
    atomic_add (&lockvar, 10000);  
    while (lockvar != 10000);  
}
```

write_release:

```
atomic_add (&lockvar, 1);
```

Readers-writer lock

- **atomic_sub**: atomically subtracts the input value from the input memory location and return the updated value
- **atomic_add**: atomically adds the input value to the input memory location and return the updated value

Problem with reader-writer lock

- Starvation
- Both readers and writer execute atomic instruction on a shared lock