



Integrity of file system in Linux

- cryptographic hash function
- Merkle tree



Shared memory IPC in xv6

- Implement IPC between parent and child
- An application can do at most one fork
- The parent process is always a sender
- The child is always a receiver

Shared memory IPC in xv6

- Implement a system call to initiate shared memory IPC
- Map a single page (shared page) in the process address space for IPC
- Always use a fixed address for the shared page
 - e.g., KERNBASE – 4096
- In fork system call, use the same physical page at (KERNBASE-4096) as the parent process

Shared memory IPC in x86

- Implement send and receive routine in the user library (**ulib.c**)
- The send and receive routine can assume that the shared page is present at (KERNBASE-4096)
- Implement a circular queue in send and receive

User mode library (**ulib.c**)

```
struct queue {  
    int head, tail;  
    char buf[4096-8];  
};  
  
// declare a global variable in ulib.c  
struct queue *ipc_queue = (struct queue*)(KERNBASE-4096);
```

- Use **ipc_queue** in send and receive implementation

User mode library

- add `send` and `receive` routines to `ulib.c`
- declare the prototype of `send` and `receive` routines in `user.h`
- add a test case similar to `memtest1` (from assignment 3) that uses the `send` and `receive` interfaces
 - search for `memtest1` in the `Makefile` of the assignment3 folder

System call handler

- Add a new system call in `syscall.c`, `syscall.h`, `sysproc.c`, `user.h`, `usys.S`
- define a new system call vector in `syscall.h`
- add a handler corresponding to vector in `syscall.c`
- define the handler in `sysproc.c`
- define the user mode system call routine in `usys.S`
- declare the prototype of user mode routine in `user.h`

Scheduling

- FIFO
- round-robin
- priority scheduling
- priority donation

How a mature OS implement scheduling?

- Multilevel feedback queue scheduler
 - maintains multiple queues
 - each queue has a different priority
 - I/O intensive and interactive workloads are given highest priority
 - compute-intensive applications have lower priority
 - priority of low priority process is increased over time if it is not getting scheduled

Multiple queue feedback scheduling



How to identify an I/O intensive application

- An I/O intensive workload is likely to yield early
- A high priority thread is given smaller time quantum
- If a thread yields before exhausting the time quantum then the application is considered as I/O intensive
- If a thread exhausts its time quantum, then it is moved to the lower priority queue with larger time quantum

Multilevel feedback queue scheduling

- Parameters used for multilevel feedback queue scheduler
 - The number of queues
 - The scheduling algorithm for each queue
 - When to move a process to a lower priority queue
 - When to move a process to a higher priority queue
 - Which queue should be assigned initially

Multilevel feedback queue scheduling (example)

- Number of queues are 8
- $\text{priority}(\text{queue}[i]) > \text{priority}(\text{queue}[j])$, where $i < j$
- $\text{time_quantum}(\text{queue}[\underline{i+1}]) = 2 * \text{time_quantum}(\text{queue}[i])$
- A new thread always added to $\text{queue}[0]$

0	1 ms
1	2 ms
2	4 ms
3	8 ms
4	16 ms

Multilevel feedback queue scheduling (example)

- If a thread in `queue[i]` yields before `quantum(queue[i])`, then it is moved to `queue[i-1]` (if exists)
- If a thread in `queue[i]` yields after `quantum(queue[i])`, then it is moved to `queue[i+1]` (if exists)
- If a thread in `queue[i]` was not scheduled in last “x” time quantum move it to `queue[i-1]` (if exists)
- Individual queue implements round-robin scheduling

Multiple-Processor scheduling

- Processor affinity
 - most multiprocessor OS try to schedule a process on the same processor every time
 - to take the full advantage of per-CPU cache
 - However, this is not possible all the time (for example, consider a case when a CPU is idle, but a process is waiting to be scheduled on other CPU)
- Linux provides system calls to set the processor affinity to a set of cores
 - In this case, a process or thread will always be scheduled on a set of CPUs

What is inside an OS?

- scheduler
- exception handlers
- system call handlers
- device drivers
- file system
- network stack (TCP/IP)
- accelerators (GPU)
- and so on

How does an OS provide process isolation?


- Memory isolation
 - Using page tables
- Disallow execution of privileged instructions
 - Using protection rings

Do we need protection rings if the compiler is trusted?

- Let us assume that the compiler never generates a privileged instruction

Do we need protection rings if the compiler is trusted?

- Some instructions have different meanings in different protection rings
 - pushf and popf instructions save and restore the eflags registers on/from the stack
 - in ring 0, popf restores the interrupt flag from the stack



Do we need protection rings if the compiler is trusted?

- Can you execute arbitrary byte code in a C program

Do we need protection rings if the compiler is trusted?

```
void foo () {
```

```
...
```

```
}
```

```
int main () {
```

```
void (*fnptr()) = foo;
```

```
...
```

```
fnptr ();
```

```
return 0;
```


```
}
```

foo:
→ b8 fa 00 00 00 mov \$0xfa, %eax
 ↑
 fa cli

→ fnptr++;
ptr = x 0x1000
jmp ~~ptr~~ 0x1000

Do we need protection rings if the compiler is trusted?

```
foo () {  
    char str[128];  
    scanf ("%s", str);  
    ...  
}
```



Next class

- How can we build an OS which offers better security?