



# End semester examination

- Open book examination
- You have to bring the xv6 code listing
- Electronic devices are not allowed

# Syllabus

- Everything that we have discussed so far
- demand paging and swapping
  - paging assignment, replacement policies, page tables, TLB, memory-mapped files, etc.
- locks
  - spinlocks, ticket spinlocks, readers-writer lock, big-reader lock, read-copy update, Peterson's solution, semaphores, etc.
- Filesystem
  - filesystem structure in xv6, buffer cache, in-memory inode, disk inode, synchronization, crash recovery, logging, etc.

# Syllabus

- CPU cache, cache coherence, weak memory model, memory barriers, etc.
- Security in OS, temporal safety, scheduling policies
- calling conventions
- Anything else that is missing here but has been discussed in the class

# How does an OS provide process isolation?

- Memory isolation
  - Using page tables
- Disallow execution of privileged instructions
  - Using protection rings

# Do we need protection rings if the compiler is trusted?

- Assuming that a “C” compiler never generates a privileged instruction, can we execute the program in ring-0
  - no, because the application can modify a function pointer or the return address on the stack to jump in the middle of an instruction

# Do we need protection rings if the compiler is trusted?

- Can we execute a Java program in ring-0

# Can we overwrite function pointers in Java


- Functions are the member of the class
  - No concept of pointer arithmetic in Java
- Typecasts of objects are always safe in Java



# Typecast in Java

```
class Type1 {  
    int field1;  
    Type1 ();  
    void foo ();  
    void bar ();  
}
```

```
class Type2 {  
    int field1;  
    int field2;  
    int field3;  
    Type2 ();  
}
```



```
Type1 obj1 = new Type1();
```

```
Type2 obj2 = (Type2)obj1; x not allowed in Java
```

```
obj2.field3 = 0x8210023;
```

```
obj1.foo();    // obj1.foo had been probably overwritten by field3
```

# How about use-after-free

```
class Type1 {  
    int field1;  
    Type1 ();  
    void foo ();  
    void bar ();  
}
```

```
class Type2 {  
    int field1;  
    int field2;  
    int field3; ✓  
    Type2 ();  
}
```

```
Type1 obj1 = new Type1();    0x1234
```

```
free (obj1);    ✗ no free in Java.
```

```
Type2 obj2 = new Type2();    — 0x1234
```

```
obj1.foo();    // obj1.foo had been probably overwritten by field3
```

# No free in Java

- How does Java detect free objects

- mark and sweep garbage collection
- arbitrary latencies

```
struct list {  
    node * node = null;  
    struct list * ptr = a list;  
    node = ptr;  
    ptr = NULL;  
};
```

node

# How about overwriting return address

```
int[] array = new int[32];  
array[34] = middle_of_some_instruction;
```

Not possible!

Java throws a runtime exception on out of bound array access.

# We can't execute privilege instruction in Java because

- No function pointer corruption
- No out of bound array access

# How about memory isolation

- Java does not provide virtual address abstraction
- applications can access objects allocated through the memory allocation primitives
- If the allocator is functionally correct, a Java application can never access an arbitrary memory location

# Memory isolation in Java

- No hardware support is needed for memory isolation
- Paging can be disabled
  - No TLB pressure, reloading of page tables, no address translation
- All applications are Java applications
- OS is also a Java application

# System calls in Java OS

- use context switch API
- allocate objects in the OS address space
- copy arguments to the OS address space
- switch to the OS
  - no **cr3** reload, no ring transition, direct jump to the **system call handler**



# IPC in Java OS

- allocate objects in the target process address space
- copy arguments to the target process address space
- switch to the target process
  - no **cr3** reload, no ring transition, jump to the **receive** routine

# Device drivers can still access OS objects

- Can we do better?

# Device driver isolation

- Each device driver is a sperate Java process
- A device driver can access OS objects by doing system calls similar to other processes

# Java OS

- Despite all these advantages why do operating systems still prefer C
- Java is not an efficient language
  - garbage collection can cause arbitrary latencies when memory utilization is high
- An efficient and safe programming language is an important and challenging problem to solve



# Can we design a relatively secure OS in C

- Problem
  - kernel subsystems, e.g., file system, network stack, device drivers, page fault handler, process manager, etc., can access all the memory
  - all of these together is a huge code base
  - if any of these components are vulnerable to exploits, the whole system can be compromised

# Microkernel

- Each subsystem runs as a separate user-mode process
- OS does minimal work
  - address space
  - threads
  - IPC (send and receive)

# Microkernel



# Pros

- fault tolerance
  - OS can restart the file system process after a null pointer dereference
- security
  - A security bug in the file system won't affect other kernel services or leak data of kernel



# Cons

- System calls are slow
  - page table switch
  - additional two transitions between different privilege levels

# Interesting problems

- OS security
- Cloud security
- Scalable OS
- OS for heterogenous architectures
- Distributed OS
- and many more.