

PC Architecture and Processor Setup

Outline

- PC architecture
- x86 instruction set

PC architecture

- A full PC has:
 - one or more x86 CPUs, each containing:
 - integer registers (can you name them?) and execution unit
 - floating-point/vector registers and execution unit(s)
 - memory management unit (MMU)
 - multiprocessor/multicore: local interrupt controller (APIC)
 - memory
 - disk (IDE, SCSI, USB)
 - keyboard
 - display
 - other resources: BIOS ROM, clock, ...
- We will start with the original 16-bit 8086 CPU (1978)
- CPU runs instructions:

```
for (;;) {  
    run next instruction  
}
```

- Draw figure with common bus, I/O, and CPU. The CPU has registers, cache, etc.
- Draw figure showing EIP and how it gets incremented automatically after executing each instruction.
- Needs work space: registers
 - four 16-bit data registers: AX, BX, CX, DX
 - each in two 8-bit halves, e.g. AH and AL
 - very fast, very few
- More work space: memory
 - CPU sends out address on address lines (wires, one bit per wire)
 - Data comes back on data lines
 - *or* data is written to data lines
- Add address registers: pointers into memory
 - SP - stack pointer
 - BP - frame base pointer
 - SI - source index
 - DI - destination index
- Instructions are in memory too!
 - IP - instruction pointer (PC on PDP-11, everything else)
 - increment after running each instruction
 - can be modified by CALL, RET, JMP, conditional jumps
- Want conditional jumps
 - FLAGS - various condition codes
 - whether last arithmetic operation overflowed
 - ... was positive/negative
 - ... was [not] zero
 - ... carry/borrow on add/subtract
 - ... etc.
 - whether interrupts are enabled
 - direction of data copy instructions
 - JP, JN, J[N]Z, J[N]C, J[N]O ...
- What if we want to use more than 2^{16} bytes of memory?
 - 8086 has 20-bit physical addresses, can have 1 Meg RAM
 - the extra four bits usually come from a 16-bit "segment register":
 - CS - code segment, for fetches via IP
 - SS - stack segment, for load/store via SP and BP
 - DS - data segment, for load/store via other registers
 - ES - another data segment, destination for string operations
 - virtual to physical translation: $pa = va + seg * 16$
 - e.g. set CS = 4096 to execute starting at 65536

- tricky: can't use the 16-bit address of a stack variable as a pointer
- a *far pointer* includes full segment:offset (16 + 16 bits)
- tricky: pointer arithmetic and array indexing across segment boundaries
- But 8086's 16-bit addresses and data were still painfully small, so 80386 added support for 32-bit data and addresses (1985)
 - boots in 16-bit mode, bootasm.S switches to 32-bit mode
 - registers are 32 bits wide, called EAX rather than AX
 - operands and addresses that were 16-bit became 32-bit in 32-bit mode, e.g. ADD does 32-bit arithmetic
 - prefixes 0x66/0x67 toggle between 16-bit and 32-bit operands and addresses: in 32-bit mode, MOVW is expressed as 0x66 MOVW
 - the .code32 in bootasm.S tells assembler to generate 0x66 for e.g. MOVW
 - 80386 also changed segments and added paged memory...
- Example instruction encoding


```
b8 cd ab 16-bit CPU, AX <- 0xabcd
b8 34 12 cd ab 32-bit CPU, EAX <- 0xabcd1234
66 b8 cd ab 32-bit CPU, AX <- 0xabcd
```
- ...and even 32 bits eventually wasn't enough, so AMD added support for 64-bit data addresses (1999)
 - registers are 64 bits wide, called RAX, RBX, etc.
 - 8 more general-purpose registers: R8 thru R15
 - boot: *still* go thru 16-bit and 32-bit modes on the way!

x86 Instruction Set

- Intel syntax: `op dst, src` (Intel manuals!)
- AT&T (gcc/gas) syntax: `op src, dst` (labs, xv6)
 - uses b, w, l suffix on instructions to specify size of operands
- Operands are registers, constant, memory via register, memory via constant
- Examples:

<u>AT&T syntax</u>	<u>"C"-ish equivalent</u>	
<code>movl %eax, %edx</code>	<code>edx = eax;</code>	<i>register mode</i>
<code>movl \$0x123, %edx</code>	<code>edx = 0x123;</code>	<i>immediate</i>
<code>movl 0x123, %edx</code>	<code>edx = *(int32_t*)0x123;</code>	<i>direct</i>
<code>movl (%ebx), %edx</code>	<code>edx = *(int32_t*)ebx;</code>	<i>indirect</i>
<code>movl 4(%ebx), %edx</code>	<code>edx = *(int32_t*)(ebx+4);</code>	<i>displaced</i>
- Instruction classes
 - data movement: MOV, PUSH, POP, ...
 - arithmetic: TEST, SHL, ADD, AND, ...
 - i/o: IN, OUT, ...
 - control: JMP, JZ, JNZ, CALL, RET
 - string: REP MOVSB, ...
 - system: IRET, INT