

Creating kernel's page table, Initializing physical memory allocator, Running Processes in Action

Creating kernel's page table

Initializing physical memory allocator

- Remember that setupkvm allocated phys mem for page directory.

How does memory allocation work?

sheet 27

- Physical memory allocator interface:

allocates a page at a time

va = kalloc()

kfree(va)

always allocates from phys mem above where kernel was loaded

so pa is va - KERNBASE

- what does kernel use phys mem allocator for?

page table pages

kernel data structures (pipe buffers, stacks, &c)

user memory

- How does the allocator work?

data structure: a list of free physical pages

allocation = remove first entry from list

free = add page to head of list

list's "next" pointers stored in first 4 bytes of each free page

that memory is available since they are free

- Allocator depends on phys pages having virtual addresses!

since it must write them to manipulate free list

just like VM code needed to write page table pages

thus xv6 maps all phys pages into kernel address space

which burns up a lot of virtual address space, limits max user mem

other arrangements are possible

- where does allocator get initial pool of free physical memory?

kinit1

called by main

end is first address beyond the end of the kernel *as a virtual address*

memory beyond that is unused

- freerange:

PGROUNDUP since newend may not be page-aligned

- Q: why must allocated pages have page-aligned addresses?

- kinit1 assumes physical memory goes up to 4MB. Hence $P2V(4 * 1024 * 1024)$.

- Uses this space for allocation of initial data structures (e.g., first page table kpgdir)

- kinit2 assumes physical memory goes up to PHYSTOP (lame)

- kfree each page

usually called on previously-allocated memory

kinit is abusing kfree a little bit

- kfree

linked list

note cast at 2827

2828 is where we depend on phys mem being mapped at a virt addr

- kalloc takes the first element of the free list

- Q: how to allocate mem for a data structure (e.g. array) > 4096 bytes?

Processes in Action

process execution states:

diagram: user/kernel, process mem, kernel thread, kern stack, pagetable

process might be executing in user space

with cr3 pointing to its page table

user mode, so can use PTE_U PTEs, < 0x80000000

or might be in a system call in the kernel

e.g. open() finding a file on disk

process's "kernel thread"

kernel mode, so can use non-PTE_U PTEs

using kernel stack

or not currently executing

xv6 has two kinds of transitions

- trap + return: user->kernel, kernel->user
 - system calls, interrupts, divide-by-zero, &c
 - hw+sw saves user registers on process's kernel stack
 - save user process state ... run in kernel ... restore state
- process switch: between kernel threads
 - one process is waiting for input, run another
 - or time-slicing between compute-bound processes
 - save p1's kernel-thread state ... restore p2's kernel-thread state

Q: why per-process kernel stack?

what would go wrong if syscall used a single global stack?

how does xv6 store process state?

- struct proc sheet 20
- kernel proc[] table has an entry for each process
- discuss pgdir, kstack, state, pid fields.
- then discuss ofile field (fd table)

discuss how any system call works.

discuss how fork system call works.

discuss how external timer interrupt works.

discuss context switching

discuss proc.tf (trapframe)

discuss proc.context (for context switch)