

Process Structure, Switching

process execution states:

- diagram: user/kernel, process mem, kernel thread, kern stack, pagetable
- process might be executing in user space
 - with cr3 pointing to its page table
 - user mode, so can use PTE_U PTEs, < 0x80000000
- or might be in a system call in the kernel
 - e.g. open() finding a file on disk
 - process's "kernel thread"
 - kernel mode, so can use non-PTE_U PTEs
 - using kernel stack
- or not currently executing

xv6 has two kinds of transitions

- trap + return: user->kernel, kernel->user
 - system calls, interrupts, divide-by-zero, &c
 - hw+sw saves user registers on process's kernel stack
 - save user process state ... run in kernel ... restore state
- process switch: between kernel threads
 - one process is waiting for input, run another
 - or time-slicing between compute-bound processes
 - save p1's kernel-thread state ... restore p2's kernel-thread state

Q: why per-process kernel stack?

what would go wrong if syscall used a single global stack?

how system call works:

- user process uses software interrupt instruction
 - int \$0x80
- trapframe pushed on top of kstack
 - partly by hardware (eip, cs, eflags, esp, ss)
 - rest by software (first few instructions in handler are push instructions)
 - handler calls the appropriate functions (e.g., syscall functions)
- syscall arguments passed from user to kernel in registers
 - kernel function can access user register values in trapframe
- syscall return value passed in register (eax)
 - kernel function overwrites trapframe->eax with the return value
- on syscall return (through iret), the trapframe is popped from kstack
 - partly by software (the last few instructions before iret are pops)
 - partly by hardware (iret pops eip, cs, eflags, esp, ss)

Discuss how syscall arguments and return values can be passed through pointers as the kernel and the user process are living in the same address space. This makes for very efficient communication between user and kernel, and justifies the 2GB of virtual address space that the kernel eats up from the process.

In an alternate organization, the kernel could have lived in a different address space, in which case, communication would have involved copying data from one address space to another. This has the advantage of strong isolation between different components (E.g., user/kernel), but are less efficient. "Microkernels" take this approach, and are used in safety-critical applications, like embedded systems.

if process calls "yield" syscall or if an external timer interrupt is received:

- if executing in user mode (definitely true for yield)
 - hardware switches to kernel stack (per process) and pushes eip, cs, eflags, esp, ss at top
 - first few instructions push of handler push the rest of user registers
 - handler calls the appropriate functions (e.g., scheduler)
 - scheduler decides if the current process should continue running?
- if so
 - it simply returns from the scheduler
 - eventually, returning to the user mode by popping the trapframe from kstack
 - trapframe popped partly by software and partly by hardware (just as before)
- if not
 - do context-switch:
 - the scheduler switches to new process's kstack (which needs to run next)
 - old process's kstack contains all information about its trapframe, and its callchain from handler to scheduler
 - the old process kstack is linked through the its pcb (process control block)
 - the saved kstacks are in a state such that switching to them starts running the process exactly from where it was preempted
 - this means that they contain the eip value from where to resume (in kernel)
 - they also contain the values of other registers, so execution can resume as though it was never interrupted.
 - typically, on resumption of the execution, the new process will return from the scheduler, eventually returning to user mode by popping the (saved) trapframe from kstack
 - trapframe popped partly by software and partly by hardware (just as before)

```

if executing in kernel mode
  hardware does not need to switch stacks. it simply pushes eip, cs, eflags
  handler saves other registers and calls scheduler (as before)
  scheduler either returns or context-switches.
  if it returns, it simply resumes execution in kernel mode from where it
  was preempted
  if it context-switches, it saves kstack, so that future scheduling of
  this process can resume from where it was preempted (by returning from
  the scheduler, as though it was never switched-out).

```

Thus, a saved kstack contains a trapframe at its top
 somewhere it also contains a "context" frame, which is used
 to save resume kernel-mode execution
 A kstack may contain more than one trapframes
 - for example, if a timer interrupt was received in the middle of kernel execution

```

kernel stack diagram (saved):
  top ->
          esp, ss
          eip, cs
          ...
          gs fs es ds
p->tf ->  edi & 7 other registers
          ---

          ... (information about function call chain inside kernel)
          ... (e.g., return address, arguments, local vars, etc)

          ---
          eip
          ebp
          ebx
          esi
p->context -> edi
          ---

p->kstack -> ...
            (bottom)

```

Overheads of Timer Interrupt Handling

Let's say, you have written a fancy program that really cares about its speed. Also, assume that when this program is running on the system, it is the only program on the system. So, whenever a timer interrupt is received, execution switches to kernel mode, where the scheduler is called, only to realize that it is the only process that is running, and the scheduler returns without switching the kernel stack (i.e., no context switch). So, you are naturally worried about the unnecessary overhead of handling the timer interrupt periodically, because it serves no purpose (it merely returns to executing the same process).

Let's look at the overhead of timer interrupt handling:

Typical timer interrupt frequencies: 10-100 Hz (i.e., every 10-100ms)

Approximate time to execute one instruction: 1-100 nanosecond (1-100ns)

Approximate number of instructions executed in one timer interval:
 $10\text{ms}/1\text{ns} = 10\text{ms}/10\text{ns} = 10^6$

Approximate number of instructions executed for handling the timer interrupt:
 say 100-1000 (conservatively 10^3)

Conservative overhead: $10^3/10^6 = 0.1\%$

Hence, the overhead is insignificant. In fact, it is possible to further increase timer frequency (to say 1000Hz) without incurring noticeable overheads. However, as we will discuss later in our discussions on scheduling, timer frequency beyond a certain value is not very useful.