

Fork, Process kstack, Context Switching

Process kstack

How to determine the kstack size?

- look at max function call depth in your code
 - corollary: do not use deep function call chains
- look at max size of local variables in functions
 - corollary: do not allocate large variables on stack. allocate them on heap if needed.
- look at size of trapframe, etc.

xv6 uses a 4KB kstack (per process). Linux uses 8KB kstacks.

BUT: we said that any trap causes a trapframe to get pushed on the kstack. So if there are many external interrupts, one after another, could the kstack overflow?

Solution:

- Ensure that the kernel cannot cause any software interrupt or exception while executing a handler.
- Ensure that handlers of external interrupts (e.g., timer, disk, network, etc.) always execute with interrupts disabled.

This ensures that there can be at most two trapframes on the kstack: the first due to a software interrupt/exception, and the second due to an external interrupt received while executing in the kernel. This allows you to calculate an upper-bound on the stack size.

What happens if an interrupt is received while the CPU was executing with disabled interrupts? The interrupt is ignored:

- Typically, the interrupting hardware expects an acknowledgement from the CPU that it received the interrupt. If it does not receive an acknowledgement, it retries the interrupt
- Also, the CPU has a buffer (of size 2 say) where it stores pending interrupts (i.e., interrupts that were received but not serviced). As soon as the CPU re-enables interrupts, the pending interrupts are delivered to it.

How does an OS keep time? One common way is to count the number of timer interrupts received.

Fork

how fork system call works:

```
the fork function creates a new PCB and a new kstack (child's)
it copies the trapframe from the current kstack to the child's
it changes the return value (eax) in the trapframe for the child
it also initializes a few more entries on the child's kstack
    so that it can control what are the first few instructions that run
    when the child gets scheduled.
it adds the PCB to the list of schedulable processes, and returns.
```

There is one kstack per CPU, which is used by the scheduler thread (the first thread to run on that CPU). This kstack is not associated with any process. Let's call this the scheduler's kstack.

At bootup, the CPU initializes itself to run with its scheduler kstack. Now it will allocate the first process (including its kstack), initialize it, and switch from the scheduler's kstack to the new process's kstack. The new process will get to run, and will likely call exec/fork to create more processes.

Each time a process wants to yield (either voluntarily or preemptively), it switches to the scheduler's kstack. The scheduler then finds a process to run, and switches to its kstack.

The swtch function on sheet 27 switches kstacks. Let's look at it.

```
It saves the current registers on stack
It then saves the current esp into its first argument (struct context **old)
It loads the new esp from its second argument (struct context *new)
```

The kstack of every suspended process looks like it has just been switched out from inside the swtch function.

Similarly, when a process is running on the CPU, the kstack of the scheduler (stored in global variable cpu->scheduler) also looks like

it has just been switched out using the switch function.

switch sheet 26

- save current thread's registers and stack
- load new threads stack and registers
- saves current registers on current stack, struct context, sheet 20
- expects new thread's stack to have registers in that format
- stack diagram:
 - eip *****
 - ebp
 - ebx
 - esi
 - edi
- Q: why these registers?
 - callee saved, might have caller's live variables
 - same format as struct context

switch sheet 26

- save current thread's registers and stack
- load new threads stack and registers
- saves current registers on current stack, struct context, sheet 20
- expects new thread's stack to have registers in that format
- stack diagram:
 - eip *****
 - ebp
 - ebx
 - esi
 - edi
- Q: why these registers?
 - callee saved, might have caller's live variables
 - same format as struct context

Process structure

how does xv6 store process state?

- struct proc sheet 20
- kernel proc[] table has an entry for each process
- discuss pgdir, kstack, state, pid fields.
- then discuss ofile field (fd table)

discuss alltraps and trapret on sheet 30

- in the course of a regular trap, trapret gets pushed to stack by the call instruction.
- in the case of a newly forked process (or the first process), the stack is initialized to contain trapret just below the trap frame.
- Why do we save all registers (and not just callee-saved regs) here?