

Process Creation

now let's talk about creating first process and its address space

Q: you'd expect there to be an array or something of pointers to the process's user memory. where is it? how does xv6 know what memory a process is using?

ordinarily p->pgdir and phys mem contents created by fork()
we will fake them for first process

ordinarily p->tf and p->context created by syscall/switch
we will fake them for first process

main calls userinit

userinit sheet 22
only called for first process
other processes created by fork
mimics fork+exec
create a normal-looking process
ordinary scheduler will run it
needs to fill in all struct proc entries

allocproc sheet 22
used by both fork and userinit
kernel stack setup:
trapframe w/ "saved user registers"
for us, initial user registers
eax, eip, esp, &c
trapret !
context w/ "saved kernel thread registers"
for us, initial kernel thread registers
eip
assumes that execution will resume in a special
assembly function called switch (sheet 27) at line 2722,
where there is code to pop the registers
from stack and execute the return instruction
Q: where will new kernel thread start executing?
doesn't set up trapframe b/c ordinarily copied
from parent by fork, which calls allocproc
but fork and userinit both always start thread in forkret

trapframe sheet 06
context sheet 20

kernel stack diagram:
top ->
esp, ss
eip, cs
...
gs fs es ds
p->tf -> edi & 7 other registers

trapret

eip = forkret
ebp
ebx
esi
p->context -> edi

p->kstack -> ...

Q: any guesses why there are *two* saved EIPs?

back to userinit
we know setupkvm -- only fills in kernel mappings
this is a new page table for the new process
not using it yet, will switch when new kernel thread starts
call inituvm w/ ptr to new process's user instructions

initcode.S sheet 77 : becomes `_binary_initcode_start`
user program
`exec("/init", args)`

init.c sheet 78 :
initializes fds 0, 1, 2
forks shell and waits for it to exit.
if shell exits, init exits too.

```

inituvm sheet 18
    we know kalloc and mappages(pgdir, va, sz, pa)
    initcode is tiny, fits in one page
    diagram: new mapping

Q: new page is mapped at va=0; could inituvm call memmove(0, init, sz)?

back to userinit sheet 22
    tf->esp -- user stack at top of page
    tf->eip=0 -- first instruction at bottom of page

main calls scheduler() sheet 12

scheduler sheet 24
    no longer initialization: kernel now fully running
    whenever process gives up CPU -> scheduler
    so kernel runs scheduler a lot
    look for a process that wants to run, run it
    p->state: SLEEPING, RUNNABLE, RUNNING
    scheduler looks for RUNNABLE

switchuvm sheet 17
    tell h/w to use p->stack if re-enters kernel
    sys call or interrupt
    load %cr3

let's watch switch to new process's page table:
    (gdb) break switchuvm
    (gdb) x/5i 0
0x0:      Cannot access memory at address 0x0
next past load %cr3
    (gdb) x/5i 0
same as initcode sheet 75
but we are still in the kernel, in scheduler

back to scheduler sheet 24
    mark RUNNING so no other CPU runs it
    now switch to new process's kernel stack, registers, EIP
    swtch(place to save current ESP, previously saved ESP to switch to)

let's watch:
    (gdb) break swtch
    si until esp switch...
    (gdb) x/6x $esp
    si past esp switch
    (gdb) x/6x $esp
    after: 4 regs, forkret, trapret

step into forkret sheet 24
    just returns
    allocproc set up stack to have it return to trapret
    watch out:
        release and initlog cause interrupts
        so hack source to set first=0 and pushcli
        si for iret &c -- si leaves interrupts off
    next into trapret

look at trapframe sheet 06
    (gdb) x/19x $esp
    0x0 0x23 are eip:cs
    0x1000 0x2b are esp:ss

trapret sheet 29
    pops trapret registers from stack, mostly zero
    popal pops 8 general-purpose registers
    iret pops ESP, EIP, clears supervisor flag
    x/5x $esp
    now we are executing at address 0x0 in initcode sheet 75

what does initcode do?
    traps back into the kernel to make exec() system call

```